

Data Reduction for Efficient Probit Regression

Christian Peters

September 27, 2021

Contents

1	Introduction	1
2	The Probit Model	2
2.1	Introduction as a Latent Variable Model	2
2.2	A Special Case of the Generalized Linear Model	4
2.3	Parameter Estimation	5
2.3.1	Finding the Maximum Likelihood Estimate	6
2.4	The Bayesian Perspective	9
2.4.1	Prior and Posterior Distributions	9
2.4.2	Gibbs Sampling in the Probit Model	10
3	Coresets and Sensitivity Sampling	12
3.1	Do small coresets always exist?	13
3.2	The Sensitivity Framework	19
3.3	Constructing the Coreset	23
3.3.1	A closer examination of the probit loss	23
3.3.2	Finding the sensitivity bounds	25
3.3.3	Bounding the VC-dimension	29
3.3.4	A simple two-pass algorithm for small coresets	31
4	Data Streams	33
5	Experiments	33
6	Concluding Remarks	33
7	Notes	33
7.1	VC Dimension	34
7.2	New idea for VC dimension proof	36
7.3	Online Leverage Scores	37
	References	38

1 Introduction

Content.

2 The Probit Model

The probit model is a special case of the generalized linear model (GLM) described in [McCullagh and Nelder, 1989]. It is a statistical method for analyzing binary datasets, which we introduce in the following definition.

Definition 1 (Dataset). *Let $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$ be a set containing $n \in \mathbb{N}$ pairs of observations $x_i \in \mathbb{R}^d$, $y_i \in \{0, 1\}$. We call \mathcal{D} a d -dimensional (binary) dataset.*

We can use this definition of a dataset (we will omit the term binary from now on since we will only be dealing with binary datasets in this work), to describe a whole range of possible scenarios that can be subjected to statistical analysis. For example, the x_i could represent some information of a patient, such as blood pressure or weight, and the y_i could indicate the presence or the absence of a heart disease.

In situations like this, we are often interested in modeling the relationship between the explanatory quantities x_i and the outcomes y_i . We need models, that can help us to answer questions about the data such as "Which factors increase/decrease the risk of suffering from a heart disease?", or "How likely is it, that a given patient will suffer from a heart disease?". The probit model is one of many approaches to model such a relationship in a probabilistic manner. It is described in detail in references like [McCullagh and Nelder, 1989], [Agresti, 2015] or [Fahrmeir et al., 2013].

We will outline the core assumptions of the probit model below, but instead of directly starting with its GLM formulation, we introduce it as a so-called latent variable model, which enables us to naturally arrive not only at its GLM specification, but also at a powerful sampling algorithm that enables us to efficiently apply the probit model in the realm of Bayesian data analysis.

2.1 Introduction as a Latent Variable Model

When using a probit model to analyze a d -dimensional dataset $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$, we implicitly make a set of assumptions about how the data was generated. Since it is reasonable to assume, that there is a degree of randomness involved in the data generating process, we model the y_i as realizations of independent random variables Y_i , which is the first assumption of the probit model.

The second assumption is that there is a hidden random quantity Y_i^* that is associated with each Y_i such that it directly determines its outcome:

$$Y_i = \begin{cases} 1, & \text{if } Y_i^* > 0 \\ 0, & \text{if } Y_i^* \leq 0 \end{cases} \quad (1)$$

The Y_i^* are also assumed to be independent from each other and, as already noted, unobservable, which is the reason why the Y_i^* are also called latent variables and why the probit model can also be thought of as a latent variable model.

The third and final assumption of the probit model defines the distribution of the Y_i^* and its part of the relationship between the non-random explanatory quantities x_i

and the outcomes y_i . In order to describe this relationship more concisely, we put all the observations x_i inside of a matrix $X \in \mathbb{R}^{n \times d}$ in such a way, that the i -th row of X corresponds to x_i . In the literature, this matrix X is often called the *model matrix* (see for example [Agresti, 2015]). We do the same with the Y_i^* and put them in a random vector Y^* as well, such that Y_i^* constitutes the i -th element of Y^* .

We are now ready for the third assumption of the probit model: The explanatory variables x_i influence Y_i^* in the form of a classical linear model:

$$Y^* = X\beta + \epsilon, \quad \epsilon \sim \mathcal{N}(0, \sigma^2 I), \quad (2)$$

where $\beta \in \mathbb{R}^d$ is the parameter vector of the linear model, ϵ is a normal distributed vector with independent components of mean zero and variance σ^2 , and $I \in \mathbb{R}^{n \times n}$ is the $n \times n$ identity matrix. It follows directly that Y^* is also normal distributed: $Y^* \sim \mathcal{N}(X\beta, \sigma^2 I)$.

These three assumptions are already a complete specification of the probit model and are summarized in the following definition as a brief recapitulation:

Definition 2 (Probit Model). *A d -dimensional binary dataset $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$ with model matrix $X \in \mathbb{R}^{n \times d}$ was generated by a probit model with parameters $\beta \in \mathbb{R}^d$ and $\sigma \in \mathbb{R}_{>0}$, if the following three assumptions are true:*

1. *The observations y_1, \dots, y_n are realizations of independent binary random variables Y_1, \dots, Y_n .*
2. *The outcomes of Y_1, \dots, Y_n are determined by hidden continuous random variables Y_1^*, \dots, Y_n^* by thresholding: If $Y_i^* > 0$, then $Y_i = 1$, and if $Y_i^* \leq 0$, then $Y_i = 0$.*
3. *The vector of hidden variables Y^* follows a multivariate normal distribution: $Y^* \sim \mathcal{N}(X\beta, \sigma^2 I)$, where $\beta \in \mathbb{R}^d$ and $\sigma \in \mathbb{R}_{>0}$ are the model parameters.*

Based on this definition, it is straight forward to determine the distribution of the response variables Y_i . We can calculate the probability $P(Y_i = 1)$ like this:

$$P(Y_i = 1) = P(Y_i^* > 0) = 1 - P(Y_i^* \leq 0) = 1 - P\left(\frac{Y_i^* - x_i^T \beta}{\sigma} \leq -\frac{x_i^T \beta}{\sigma}\right) = \Phi\left(\frac{x_i^T \beta}{\sigma}\right),$$

where $\Phi(\cdot)$ is the cumulative distribution function of the standard normal distribution:

$$\Phi(x) = \int_{-\infty}^x \frac{1}{\sqrt{2\pi}} e^{-\frac{1}{2}z^2} dz.$$

The result $P(Y_i = 1) = \Phi\left(\frac{x_i^T \beta}{\sigma}\right)$ leads us to an interesting observation: Both parameters β and σ are unknown model parameters and every value of σ can be compensated by a corresponding scaling of β . This means that, because we can't observe the hidden variables Y_i^* , it is impossible to determine which β and which σ generated the data without any prior knowledge. We can only draw conclusions with regard to the scaled parameter $\frac{1}{\sigma}\beta$. In this situation, we say that β and σ are *not identifiable*.

For this reason, in literature like [Fahrmeir et al., 2013] or [Agresti, 2015], it is often argued, that without the loss of generality, we can assume that $\sigma = 1$ and arrive at

$$P(Y_i = 1) = \Phi(x_i^T \beta).$$

Conversely, since Y_i is binary, it follows that

$$P(Y_i = 0) = 1 - P(Y_i = 1) = 1 - \Phi(x_i^T \beta) = \Phi(-x_i^T \beta),$$

and we arrive at the model equations:

$$Y_i \sim \text{Bin}(1, \pi_i), \quad \pi_i = \Phi(x_i^T \beta), \quad (3)$$

where $\text{Bin}(1, \pi_i)$ is a Bernoulli distribution with success probability $\pi_i = \Phi(x_i^T \beta)$.

2.2 A Special Case of the Generalized Linear Model

The final equations of the probit model that we arrived at in equation 3 are a special case of a more general model concept, the generalized linear model (GLM), that we briefly touch on below.

Generalized linear models consist of three components. The first one is the so called *random component*, a set of $n \in \mathbb{N}$ independent random variables $\{Y_i\}_{i=1}^n$. In GLMs, the distribution of these random variables is assumed to be a member of the *exponential family*, a broad family of probability distributions that encompasses the normal distribution, the binomial distribution and many others. It is characterized in more detail in [Agresti, 2015].

The second component of a GLM is the *linear predictor*. Just like in the probit model, we also assume that we are presented with some fixed observations $\{x_i \in \mathbb{R}^d\}_{i=1}^n$, that are assumed to have some explanatory power with regard to the Y_i . We thus call these observations the explanatory quantities. The linear predictor is used to relate the explanatory quantities to the distribution of the Y_i by linearly combining them as follows:

$$\eta_i = x_i^T \beta,$$

where $\eta_i \in \mathbb{R}$ denotes the linear predictor related to observation x_i and $\beta \in \mathbb{R}^d$ is the unknown parameter vector of the GLM that has to be estimated when fitting the model.

The third component of a GLM is the so called *link function*. This is a monotonic, differentiable and invertible function g that connects the linear predictor η_i to the distribution of the Y_i like this:

$$g(E[Y_i]) = \eta_i.$$

We are thus using the link function g to transform the expected value $E[Y_i]$ in such a way that it can be predicted by a linear model, hence the name *generalized linear models*.

Equivalently, we can also characterize this relationship by using the inverse function $h = g^{-1}$, also called the *response function*:

$$E[Y_i] = h(\eta_i).$$

We are now ready to establish the connection between the probit model and the generalized linear model. As we saw in equation 3, the assumptions of the probit model imply that the Y_i follow independent binomial distributions with a success probability of $\pi_i = \Phi(x_i^T \beta)$. The binomial distribution is a member of the exponential family, so we can also think of the Y_i as the random component of a GLM.

It also follows directly from the binomial distribution that $E[Y_i] = \pi_i$, thus we have from the probit model equations that $\pi_i = E[Y_i] = \Phi(x_i^T \beta)$, and equivalently $\Phi^{-1}(E[Y_i]) = x_i^T \beta$. Thus, we can think of Φ as the response function of a GLM and Φ^{-1} , the quantile function of the standard normal distribution, as the link function. The function Φ^{-1} is also known as the *probit function*, hence the name probit model.

2.3 Parameter Estimation

The parameters of generalized linear models and therefore the parameters of the probit model are usually estimated by using the *maximum likelihood method*. This method seeks to maximize the likelihood that some observed dataset $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$ was generated under the assumptions of the model, given some parameter vector $\beta \in \mathbb{R}^d$.

To make notation a little easier, we also put the outcomes y_i in a vector $y \in \{0, 1\}^n$ such that y_i is the i -th component of y . In the same way, we also put the random variables Y_i inside of a random vector Y .

In the probit model, the likelihood function is given as

$$\mathcal{L}(\beta) = P(Y = y|\beta) = \prod_{i=1}^n P(Y_i = y_i|\beta), \quad (4)$$

because the Y_i are independent. By using a little trick, we can write $P(Y_i = y_i|\beta)$ as a single expression by combining the equations $P(Y_i = 1) = \Phi(x_i^T \beta)$ and $P(Y_i = 0) = \Phi(-x_i^T \beta)$ from section 2.1 like this:

$$P(Y_i = y_i|\beta) = \Phi[(2y_i - 1)x_i^T \beta],$$

which works because $2y_i - 1 = 1$ for $y_i = 1$ and $2y_i - 1 = -1$ for $y_i = 0$. This enables us to arrive at the likelihood

$$\mathcal{L}(\beta) = \prod_{i=1}^n P(Y_i = y_i|\beta) = \prod_{i=1}^n \Phi[(2y_i - 1)x_i^T \beta] = \prod_{i=1}^n \Phi(-z_i^T \beta). \quad (5)$$

Here, we introduced the new vector $z_i = -(2y_i - 1)x_i$, which will simplify the notation later on.

The maximum likelihood estimate for β is then given by

$$\hat{\beta} \in \operatorname{argmax}_{\beta \in \mathbb{R}^d} \mathcal{L}(\beta), \quad (6)$$

and for $n \rightarrow \infty$ it holds that $E[\hat{\beta}] = \beta$ [Fahrmeir et al., 2013].

However, for finite sample sizes, the existence of $\hat{\beta}$ cannot be guaranteed and is dependent on the observed data. An overview of the conditions for the existence and uniqueness of $\hat{\beta}$ is given in [Demidenko, 2001]. In particular, there is one important condition shown in [Lesaffre and Kaufmann, 1992], that is related to the concept of linear separability, which we introduce in the following definition.

Definition 3 (Linear separability). *Let $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$ be a d -dimensional binary dataset. Let $S_0 = \{i \in [n] : y_i = 0\}$ and $S_1 = \{i \in [n] : y_i = 1\}$. If there exists a $\beta \in \mathbb{R}^d \setminus \{0\}$ such that*

$$\forall i \in S_0 : x_i^T \beta \leq 0 \quad \text{and} \quad \forall i \in S_1 : x_i^T \beta \geq 0,$$

then we call \mathcal{D} linearly separable.

Intuitively speaking, a dataset is linearly separable if there exists a hyperplane that perfectly separates the datapoints labeled with 1 from the datapoints labeled with 0. This property of a dataset is a both sufficient and necessary condition for the existence of the maximum likelihood estimate $\hat{\beta}$, as stated in the following theorem.

Theorem 1 ([Lesaffre and Kaufmann, 1992]). *Let $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$ be a d -dimensional binary dataset. The maximum likelihood estimate $\hat{\beta}$ for the parameter β of the probit model exists if and only if \mathcal{D} is not linearly separable.*

In [Haberman, 1974], it was further shown, that if the maximum likelihood estimate exists and the model matrix X has full column rank, i.e. $\text{rank}(X) = d$, then it is also unique. It now remains to explore, how the maximum likelihood optimization problem can be solved in such a case.

2.3.1 Finding the Maximum Likelihood Estimate

For the reason that the likelihood function $\mathcal{L}(\beta)$ is numerically inconvenient to maximize, the natural logarithm is often applied as a transformation to simplify the optimization problem:

$$\ell(\beta) = \ln \mathcal{L}(\beta) = \sum_{i=1}^n \ln \Phi(-z_i^T \beta). \quad (7)$$

Since we later wish to interpret ℓ as a loss function, we prefer to minimize the negative value of ℓ rather than maximizing:

$$f(\beta) = -\ell(\beta) = \sum_{i=1}^n \ln \left(\frac{1}{1 - \Phi(z_i^T \beta)} \right) = \sum_{i=1}^n g(z_i^T \beta). \quad (8)$$

Here, we define $g(x) = \ln \left(\frac{1}{1 - \Phi(x)} \right)$ and call it the *probit loss*, i.e. the loss-function that determines how much each z_i contributes to the total loss $f(\beta)$ for a given value of β .

At this point, we could already elaborate on the minimization of $f(\beta)$, but there is one more generalization that we have to make, which will later be needed when applying

the theory of data reduction to the probit model: We have to introduce positive sample weights w_1, \dots, w_n , alternatively specified by the weight vector $w \in \mathbb{R}_{>0}^n$, that give a positive weight to each datapoint in the objective function:

$$f_Z^w(\beta) = \sum_{i=1}^n w_i g(z_i^T \beta). \quad (9)$$

Here, we also introduced the subscript Z , which refers to the matrix $Z \in \mathbb{R}^{n \times d}$, where the i -th row of Z is given by z_i , but if Z and w are clear from the context, we will usually omit it and simply refer to f_Z^w by f . In the rest of this work, we will be referring to Z as the *scaled model matrix*.

The optimization of f is usually done by applying the Newton-Raphson algorithm, an iterative procedure that starts at some initial guess $\beta^{(0)}$ and successively updates it like this:

$$\beta^{(t)} = \beta^{(t-1)} - \left(\frac{\partial^2 f(\beta^{(t-1)})}{\partial \beta \partial \beta^T} \right)^{-1} \cdot \frac{\partial f(\beta^{(t-1)})}{\partial \beta}, \quad (10)$$

where $\left(\frac{\partial^2 f(\beta^{(t-1)})}{\partial \beta \partial \beta^T} \right)^{-1}$ refers to the inverse of the hessian matrix of f , evaluated at $\beta^{(t-1)}$, and $\frac{\partial f(\beta^{(t-1)})}{\partial \beta}$ refers to the gradient of f , evaluated at $\beta^{(t-1)}$. The idea behind this procedure is, broadly speaking, to approximate f locally around $\beta^{(t)}$ as a second degree taylor-polynomial and then analytically find the minimum of this polynomial. The minimum of this local polynomial approximation of f is then iteratively used as the basis for the next step of the Newton-Raphson algorithm.

It remains to find the gradient as well as the hessian matrix of f . Because f is a sum of the function g evaluated at different points, it makes sense to first determine the derivative of g . This can be accomplished by using the chain rule as follows:

$$\begin{aligned} \frac{d}{dx} g(x) &= \frac{d}{dx} \ln \left(\frac{1}{1 - \Phi(x)} \right) \\ &= (1 - \Phi(x)) \cdot \frac{d}{dx} \left(\frac{1}{1 - \Phi(x)} \right) \\ &= (1 - \Phi(x)) \cdot \frac{(-1)}{(1 - \Phi(x))^2} \cdot \frac{d}{dx} (1 - \Phi(x)) \\ &= \frac{(-1)}{1 - \Phi(x)} \cdot (-1) \cdot \phi(x) \\ &= \frac{\phi(x)}{1 - \Phi(x)}, \end{aligned} \quad (11)$$

where $\phi(x)$ is the density function of the standard normal distribution function:

$$\phi(x) = \frac{1}{\sqrt{2\pi}} e^{-\frac{1}{2}x^2}.$$

We can use this result to calculate the gradient of f :

$$\begin{aligned}
\frac{\partial}{\partial \beta} f(\beta) &= \frac{\partial}{\partial \beta} \sum_{i=1}^n w_i g(z_i^T \beta) \\
&= \sum_{i=1}^n w_i z_i g'(z_i^T \beta) \\
&= \sum_{i=1}^n w_i z_i \frac{\phi(z_i^T \beta)}{1 - \Phi(z_i^T \beta)}
\end{aligned} \tag{12}$$

Next, we need to determine the hessian matrix of f . In order to do this, we again start by finding the second derivative of g , this time using the quotient rule:

$$\begin{aligned}
\frac{d^2}{dx^2} g(x) &= \frac{d}{dx} \frac{\phi(x)}{1 - \Phi(x)} \\
&= \frac{\phi'(x)(1 - \Phi(x)) - \phi(x) \cdot (-1) \cdot \phi(x)}{(1 - \Phi(x))^2} \\
&= \frac{(-1) \cdot x \cdot \phi(x)(1 - \Phi(x)) - \phi(x) \cdot (-1) \cdot \phi(x)}{(1 - \Phi(x))^2} \\
&= \frac{[\phi(x)]^2 - x \cdot \phi(x) \cdot (1 - \Phi(x))}{(1 - \Phi(x))^2} \\
&= \left(\frac{\phi(x)}{1 - \Phi(x)} \right)^2 - x \cdot \frac{\phi(x)}{1 - \Phi(x)} \\
&= \frac{\phi(x)}{1 - \Phi(x)} \left(\frac{\phi(x)}{1 - \Phi(x)} - x \right) \\
&= g'(x) \cdot (g'(x) - x)
\end{aligned} \tag{13}$$

We can now use this result to find the hessian matrix of f :

$$\begin{aligned}
\frac{\partial^2}{\partial \beta \partial \beta^T} f(\beta) &= \sum_{i=1}^n \frac{\partial^2}{\partial \beta \partial \beta^T} w_i g(z_i^T \beta) \\
&= \sum_{i=1}^n w_i z_i z_i^T g'(z_i^T \beta) (g'(z_i^T \beta) - z_i^T \beta) \\
&= \sum_{i=1}^n w_i z_i z_i^T \frac{\phi(z_i^T \beta)}{1 - \Phi(z_i^T \beta)} \left(\frac{\phi(z_i^T \beta)}{1 - \Phi(z_i^T \beta)} - z_i^T \beta \right).
\end{aligned} \tag{14}$$

Because it can be shown, that $f(\beta)$ is a convex function [Wedderburn, 1976], and that the Newton Raphson algorithm converges to a global optimum when applied to a convex function [Nocedal and Wright, 2006], the optimization procedure converges to the maximum likelihood estimate $\hat{\beta}$ under the condition that the data is not linearly separable.

2.4 The Bayesian Perspective

The most fundamental difference between the bayesian approach to the probit model and the frequentist approach that was discussed above, is the assumption, that the model parameter β is not a fixed value, but a random variable with a probability distribution. The goal of bayesian data analysis is to draw conclusions about the distribution of the model parameter and to update these conclusions after observing more and more data.

A detailed overview of the principles of bayesian data analysis would certainly go beyond the scope of this work, but the interested reader will find a comprehensive reference in [Gelman et al., 2013]. In this section, we will merely touch on the most essential concepts, which are required in order to understand the bayesian view on the probit model.

2.4.1 Prior and Posterior Distributions

To characterize the prior uncertainty about the model parameter β , the first step of bayesian data analysis is to specify a so called *prior distribution*. In the probit model, one common choice that is also described in [Fahrmeir et al., 2013] is to assume that

$$\beta \sim \mathcal{N}(\mu_\beta, \Sigma_\beta), \quad (15)$$

i.e. β follows a normal distribution with mean $\mu_\beta \in \mathbb{R}^d$ and covariance matrix $\Sigma_\beta \in \mathbb{R}^{d \times d}$.

We can think of μ_β and Σ_β as a way to include prior knowledge into the model. If such knowledge is not present, we can choose μ_β and Σ_β in a more general fashion, perhaps we decide to set $\mu_\beta = 0$ and $\Sigma_\beta = \sigma_\beta \cdot I$ for a large value of σ_β , which would be an example of an *uninformative* prior because of the relatively unrestrictive assumptions. Alternatively, we could even go as far and also specify prior distributions on μ_β and Σ_β , which would lead us into the realm of hierarchical models (see [Gelman et al., 2013] for more details). But this would definitely go beyond the scope of this work, which is why we assume from now on that the values of μ_β and Σ_β are specified beforehand in a reasonable manner.

The next step in the process of bayesian data analysis is to determine how we should update our initial prior assumptions about β after we observed some new data represented by the dataset $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$. Ultimately, the goal is to determine the *posterior distribution* of β given the new data, represented by the probability density function $p(\beta|Y = y)$, where y is the vector of observations and Y is the random vector that we assumed to have generated these observations in the probit model.

We can find the posterior distribution by making use of the bayes rule:

$$p(\beta|Y = y) = \frac{p(Y = y|\beta)p(\beta)}{p(Y = y)}. \quad (16)$$

This relationship tells us, that in order to arrive at the posterior distribution, there are three different parts that we have to combine.

The first part is the likelihood function $p(Y = y|\beta)$, which we already dealt with in section 2.3, and the second part is the prior density function $p(\beta)$, that we assumed to be normal.

The third and most challenging part to compute is the quantity $p(Y = y)$. We can see why it is so challenging by writing it out:

$$\begin{aligned} p(Y = y) &= \int p(Y = y|\beta)p(\beta)d\beta \\ &= \int \prod_{i=1}^n \Phi(-z_i^T \beta) \frac{1}{\sqrt{(2\pi)^d \det \Sigma_\beta}} \exp\left(-\frac{1}{2}(\beta - \mu_\beta)^T \Sigma_\beta^{-1}(\beta - \mu_\beta)\right) d\beta \end{aligned} \quad (17)$$

This infinite integral over all possible values of β is impossible to solve analytically, which means that it's impossible to exactly compute the posterior distribution for the probit model. But luckily, encountering an intractable integral like this is quite common in bayesian data analysis, so there are workarounds that still allow us to analyze the posterior distribution, even though we are unable to determine it exactly.

The first consideration is, that we could also analyze the posterior distribution if we had a large enough sample of it available instead. When the sample size is big enough, the Glivenko-Cantelli theorem tells us that the empirical posterior distribution converges to the true posterior distribution [Vaart, 1998]. This allows us to analyze the posterior distribution by analyzing a large enough sample of it, but the problem of how to obtain such a sample still remains.

In practice, instead of directly sampling from $p(\beta|Y = y)$, the posterior distribution can be approximated by so called Markov chain Monte Carlo (MCMC) methods. One such method that works particularly well for the probit model is the Gibbs sampler, which is described in the next section.

2.4.2 Gibbs Sampling in the Probit Model

Gibbs sampling is an iterative tool for drawing samples from probability distributions, that was first applied in the context of bayesian inference by [Gelfand and Smith, 1990] and has been adapted to the probit model by [Albert and Chib, 1993] using the idea of *data augmentation*, which was first introduced in [Tanner and Wong, 1987]. We describe this idea in the following section, as it yields an efficient algorithm for sampling from the posterior distribution of the probit model.

Remember, that the probit model has the following components: The vector of latent variables Y^* that follows a linear model $Y^* | \beta \sim \mathcal{N}(X\beta, 1)$, where we assume that $\sigma = 1$ for reasons of identifiability (see section 2.1) and the random vector Y that produces the observed outcomes y by thresholding: If $Y_i^* > 0$, then $Y_i = 1$ and $Y_i = 0$ otherwise. We also assumed a normal prior distribution: $\beta \sim \mathcal{N}(\mu_\beta, \Sigma_\beta)$.

Now, imagine that we knew the outcomes of the latent variable vector Y^* . The conditional distribution of β given the realization y^* of the latent variables can be shown to be normal [Albert and Chib, 1993]:

$$\beta | Y^* = y^* \sim \mathcal{N}(b, B), \quad (18)$$

where $b = (\Sigma_\beta^{-1} + X^T X)^{-1}(\Sigma_\beta^{-1} \mu_\beta + X^T y^*)$ and $B = (\Sigma_\beta^{-1} + X^T X)^{-1}$. From this distribution, it is possible to sample efficiently.

The problem is, that in reality we can't observe the latent variables and therefore we don't know the realizations y^* . Here, an important finding by [Albert and Chib, 1993] comes into play: If we could observe β and see the realization $\tilde{\beta}$, then we could determine the conditional distribution of the latent variable vector Y^* :

$$Y_i^* \mid \beta = \tilde{\beta}, Y_i = y_i \sim \begin{cases} \mathcal{N}(x_i^T \tilde{\beta}, 1) \text{ truncated at the left by } 0, & \text{if } y_i = 1 \\ \mathcal{N}(x_i^T \tilde{\beta}, 1) \text{ truncated at the right by } 0, & \text{if } y_i = 0 \end{cases} \quad (19)$$

This means, that given a realization $\tilde{\beta}$ and the observed values in y , the latent variables follow a truncated normal distribution, from which it is also possible to sample efficiently.

These two observations bring us directly to the Gibbs sampling algorithm for the probit model. The first step of this procedure is to determine a starting value $\tilde{\beta}^{(0)}$. [Albert and Chib, 1993] suggest that this could for example be the maximum likelihood estimate, that we already discussed in section 2.3.

The next step of the Gibbs sampling algorithm is to use this value $\tilde{\beta}^{(0)}$ to sample a realization $y^{*(1)}$ from the latent variable vector Y^* , by using the conditional distribution in equation 19. Given $y^{*(1)}$, it is then possible to sample a new value $\tilde{\beta}^{(1)}$ from the normal distribution in equation 18, which starts a new cycle. These two sampling steps, which can both be carried out efficiently, are repeated until the desired amount of samples is reached. See Algorithm 1 for the full algorithm.

To sum up, we augmented the observed data by incorporating the hidden variables Y^* to arrive at a two-stage procedure that draws alternating samples from the conditional distributions of β and Y^* , hence the name data augmentation.

Algorithm 1: Gibbs Sampler for the Probit Model

Input: Dataset $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$ with model matrix X , prior mean $m \in \mathbb{R}^d$, prior covariance matrix $M \in \mathbb{R}^{d \times d}$, sample size $k \in \mathbb{N}$

Output: A sample β_1, \dots, β_k from the posterior distribution

```

1 Set  $B = (M^{-1} + X^T X)^{-1}$ 
2 Initialize  $\beta_0 = \hat{\beta}$ , where  $\hat{\beta}$  is the MLE for  $\beta$  computed on  $\mathcal{D}$ 
3 for  $j = 1, \dots, k$  do
4   for  $i = 1, \dots, n$  do
5     if  $y_i = 1$  then
6       Sample  $y_i^{*(j)}$  from  $\mathcal{N}(x_i^T \beta_{j-1}, 1)$  truncated at the left by 0
7     else if  $y_i = 0$  then
8       Sample  $y_i^{*(j)}$  from  $\mathcal{N}(x_i^T \beta_{j-1}, 1)$  truncated at the right by 0
9   Set  $y^{*(j)} = (y_1^{*(j)}, \dots, y_n^{*(j)})^T$ 
10  Set  $b^{(j)} = B (M^{-1} m + X^T y^{*(j)})$ 
11  Sample  $\beta_j$  from  $\mathcal{N}(b^{(j)}, B)$ 
12 return  $\beta_1, \dots, \beta_k$ 
```

3 Coresets and Sensitivity Sampling

The Newton-Raphson algorithm for optimizing the objective function of the probit model, as well as the Gibbs sampler, are reasonably efficient when the datasets are of small to moderate size. Usually, this is the case when it's possible to store the model matrix X into the main memory. But problems arise, when the datasets are getting so big, that this is no longer possible. What should we do in such a case?

One idea to deal with this issue is that we could select a smaller subset \mathcal{C} of our initial dataset \mathcal{D} , that represents the characteristics of the original data well in some sense. We hope, that when we execute the computationally expensive optimization algorithms on the smaller subset, we still get similar results as if we executed the algorithms on the original dataset. But what does it mean for a subset \mathcal{C} to be representative of \mathcal{D} ? And how could we come up with an algorithm that selects such a subset efficiently? The method of *coresets* (see for example [Munteanu and Schwiiegelshohn, 2018]) is one way to deal with these questions and we will explore it in this chapter.

So, what is a coreset? As the name suggests, we are talking about a subset $\mathcal{C} \subseteq \mathcal{D}$ of our initial dataset, which fulfills some very special requirements which ensure that the original dataset is well represented for our problem. To be more specific, we want the objective function $f(\beta)$ evaluated on the coreset to be as close to the objective function on the original dataset as possible, for all possible $\beta \in \mathbb{R}^d$. Mathematically speaking, what we are interested in is a so-called $(1 \pm \epsilon)$ approximation of the objective function on the original dataset.

To understand what is meant by that, assume for a moment that we are given a function $f(\beta)$ and an approximation $\tilde{f}(\beta)$. If $\tilde{f}(\beta)$ is a $(1 \pm \epsilon)$ approximation of $f(\beta)$, it will never deviate from $f(\beta)$ more than a factor $(1 \pm \epsilon)$, i.e. we have for all $\beta \in \mathbb{R}^d$, that:

$$(1 - \epsilon)f(\beta) \leq \tilde{f}(\beta) \leq (1 + \epsilon)f(\beta).$$

This kind of approximation would then allow us to run an optimization algorithm on $\tilde{f}(\beta)$ and guarantee that our solution is close to the optimal solution on $f(\beta)$.

As we already hinted at, a coreset is simply a subset $\mathcal{C} \subseteq \mathcal{D}$ of our original dataset, that provides us with a $(1 \pm \epsilon)$ approximation on the original loss function. We formalize this concept in the following definition.

Definition 4 (Coreset). Let $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$ be a d -dimensional dataset with scaled model matrix $Z \in \mathbb{R}^{n \times d}$, i.e. $z_i = -(2y_i - 1)x_i$ constitutes the i -th row of Z , and let $w \in \mathbb{R}_{>0}^n$ be a vector of positive sample weights. Let $\mathcal{C} \subseteq \mathcal{D}$ be a subset of \mathcal{D} of size $|\mathcal{C}| = k$ with scaled model matrix $C \in \mathbb{R}^{k \times d}$ and a vector of positive sample weights $u \in \mathbb{R}_{>0}^k$. Let $\epsilon > 0$. We call \mathcal{C} a $(1 + \epsilon)$ -coreset of \mathcal{D} for probit regression, if

$$(1 - \epsilon)f_Z^w(\beta) \leq f_C^u(\beta) \leq (1 + \epsilon)f_Z^w(\beta) \quad \forall \beta \in \mathbb{R}^d,$$

where $f_Z^w(\beta) = \sum_{i=1}^n w_i g(z_i^T \beta)$ is the weighted objective function of the probit model and $g(x) = \ln \left(\frac{1}{1 - \Phi(x)} \right)$ is the probit loss.

One thing to note here is that we do not only need to select the subset $\mathcal{C} \subseteq \mathcal{D}$, but we also have to come up with some new sample weights u . Intuitively speaking, this makes sense because when reducing the amount of data points in the objective function, which is achieved by selecting the subset \mathcal{C} , we are also naturally lowering its overall value, since g , the probit loss, is a positive function. The reweighting by u accounts for that, so we can still get a $(1 \pm \epsilon)$ approximation of the original loss function.

Let us now consider perhaps the most important aspect of this definition, which will determine the usefulness of *any* work in the domain of coresets: The coreset size $k = |\mathcal{C}|$.

It can easily be verified, that we can always come up with a coreset when $k = n$, i.e. the so-called trivial coreset, where we simply select $\mathcal{C} = \mathcal{D}$. But such a coreset doesn't help us at all with our goal of reducing the computational burden of the optimization and Gibbs sampling algorithms. Informally speaking, we want k to be small. But how small is small enough? Usually, we can consider it a success, if we can find coresets where $k \in O(\log(n))$, using the big-o notation to indicate that k is not much larger than the logarithm of the amount of data points. If, for example, we had a dataset with one billion observations, i.e. $n = 1,000,000,000$, the natural logarithm of n would equal to roughly 20 datapoints. That sounds like a decent compression, doesn't it?

In the remainder of this work, we will refer to a coreset as small, if k is roughly logarithmic in n . Our goal is to construct algorithms, which will enable us to find such small coresets in the context of probit regression.

3.1 Do small coresets always exist?

Before attempting to construct an algorithm that is able to find small coresets, we first have to investigate if such a goal is even attainable, i.e. we have to make sure that small coresets even exist.

Without imposing any restrictions on the datasets, it turns out that it is not difficult to find a counter example, i.e. to find a dataset \mathcal{D} , such that no subset $\mathcal{C} \subseteq \mathcal{D}$ of roughly logarithmic size can be a coreset. This negative result has first been proven in the context of logistic regression by the authors of [Munteanu et al., 2018], but it turns out that the problematic dataset that admits no small coresets is the same for probit regression as well.

This finding forces us to make a decision: Do we have to give up our search for small coresets because we now know that they don't always exist? Or is there still hope, perhaps by imposing some (very reasonable) restrictions on the class of datasets that we consider? Before we can turn to this discussion, we first reproduce the counter example for the general case with the following theorem.

Theorem 2. *There exists a dataset \mathcal{D} of size $|\mathcal{D}| = n$, such that any $(1 + \epsilon)$ -coreset \mathcal{C} of \mathcal{D} for probit regression has a size $k = |\mathcal{C}|$ of at least $k \in \Omega\left(\frac{n}{\log n}\right)$.*

Proof. We can construct such a dataset by showing how coresets can be used in a communication protocol for the so called INDEX game, a communication game for two players, Alice and Bob, which works like this:

Alice is given a random binary string $m \in \{0, 1\}^n$ of n bits and Bob is given an index $i \in [n]$. The objective of the game is for Alice to send a message to Bob that allows Bob to obtain the value m_i of Alice's binary string m . It was shown in [Kremer et al., 1999], that the minimum length of a message sent by Alice that still allows Bob to obtain m_i with constant probability is in $\Omega(n)$ bits. We will now see, how a coresset for probit regression can be used to encode such a message.

The first step is for Alice to convert her binary string m into a dataset \mathcal{D} as follows: For each entry m_j of her binary string where $m_j = 1$, she adds a point

$$x_j = \left(\cos \left(2\pi \frac{j}{n} \right), \sin \left(2\pi \frac{j}{n} \right), 1 \right)^T$$

to her set \mathcal{D} and labels it with $y_j = 1$, ending up with the dataset

$$\mathcal{D} = \{(x_j, 1)\}_{j \in \{i \in [n]: m_i=1\}},$$

with all points being on the unit circle.

The next step for her is to construct a $(1 + \epsilon)$ -coreset \mathcal{C} of \mathcal{D} for probit regression with sample weights $u \in \mathbb{R}_{>0}^k$ and to transmit both the coresset and the weight vector to Bob, which requires $O(\log(n))$ space for each point and weight. We will later see, how large the size $|\mathcal{C}| = k$ of this coresset must be, so that Bob can still obtain the value of m_i with constant probability.

As soon as Alice's coresset \mathcal{C} arrives at Bob, Bob can use it to obtain the value of m_i . To do this, Bob first adds two new points

$$q_1 = \left(\cos \left(2\pi \frac{i - 0.5}{n} \right), \sin \left(2\pi \frac{i - 0.5}{n} \right), 1 \right)^T$$

and

$$q_2 = \left(\cos \left(2\pi \frac{i + 0.5}{n} \right), \sin \left(2\pi \frac{i + 0.5}{n} \right), 1 \right)^T$$

to the set and labels both points with 0 (see figure 1), i.e. Bob now has the dataset

$$\mathcal{C}' = \mathcal{C} \cup \{(q_1, 0)\} \cup \{(q_2, 0)\}.$$

Next, he uses this new dataset \mathcal{C}' with scaled model matrix C' to minimize the weighted objective function $f_{C'}^u$ of the probit model, by using the Newton-Raphson optimization algorithm.

Taking a look at figure 1, it becomes evident, that Bobs points q_1 and q_2 are linearly separable from the other points if and only if Alice didn't add a point x_i , i.e. if $m_i = 0$. He can use the results of the optimization procedure to make a distinction between the two cases (which then allows him to determine the value of m_i) like this:

In the case of $m_i = 1$, Bobs points are not linearly separable from Alices original points, which means that there must occur at least one misclassification at a cost of $g(0) = \log(2)$ for the original loss function. Because Bobs dataset \mathcal{C}' allows him to

obtain a $(1 \pm \epsilon)$ -approximation of the original cost function, he can check if the Newton-Raphson algorithm converges to a cost of at least $(1 - \epsilon) \log(2)$. In this case, he knows that Alice must have added the point x_i , which means that $m_i = 1$.

Conversely, if at any point during the optimization procedure the cost function drops below $(1 - \epsilon) \log(2)$ and approaches zero, Bob knows that Alice didn't add the point x_i , because his dataset \mathcal{C}' is linearly separable. This will allow him to conclude that $m_i = 0$.

Let us now see, how large the size k of Alice's coreset must be for this protocol to work with constant probability. In [Kremer et al., 1999] it was shown, that the minimum length of a message that Alice must send in order for the protocol to work is in $\Omega(n)$ bits. Since each of the points that Alice created can be encoded in $\log(n)$ space, it follows from the lower bound that $\Omega(n) \subseteq \Omega(k \log(n))$, so k must be in $\Omega\left(\frac{n}{\log(n)}\right)$.

We can conclude, that if there existed a $(1 + \epsilon)$ -coreset of Alice's dataset \mathcal{D} for probit regression with size $k \in o\left(\frac{n}{\log(n)}\right)$, it would contradict the minimum message length of the INDEX communication game, which proves the theorem. \square




Figure 1: Bob places two points q_1 and q_2 in such a way on the unit circle, that they can be linearly separated from the other points if and only if Alice didn't place a point at x_i .

We now have an example of a dataset for which no small coresets exist, which implies that in the general case, without any restrictions, there are no guarantees that it's even possible to find a small coreset. But there is one thing that we have to note: The counter example from the INDEX proof is by no means a dataset that could ever be reasonably subjected to a probit analysis. It consists of only positive labels! How can any inference ever be possible in such a case? Further, it is easy to recognize, that the counter example is linearly separable. As we already saw in section 2.3, when estimating the model parameters, the maximum likelihood estimate only exists when the data is

not linearly separable. So yes, we found an example dataset for which no small coresets exist, but does that mean that this particular "degenerate" example is relevant to the attainment of our goal of constructing efficient data reduction algorithms for the purpose of probit regression? Since the maximum likelihood estimate doesn't even exist for this dataset, it can be doubted, to say the least.

It therefore seems reasonable to impose some restrictions on the datasets under study. Since we are exclusively dealing with probit regression, it makes sense to restrict the class of data sets to those, where a probit model can at least be properly estimated, i.e. where the data is not linearly separable and where the maximum likelihood estimate exists.

The authors of [Munteanu et al., 2018] were dealing with similar issues in the context of logistic regression, so they decided to introduce a measure, which they call μ , that describes the degree of separability of a dataset. In their work, they were able to not only use this measure to restrict the class of datasets under study, which they defined as μ -complex, but also to relate the size of their coresets directly to μ . We will go down a similar path in the search of small coresets for probit regression, so the first step for us is to slightly adapt this measure for our purposes:

Definition 5. (*μ -complexity*) Let \mathcal{D} be a d -dimensional dataset of size $|\mathcal{D}| = n$ with scaled model matrix $Z \in \mathbb{R}^{n \times d}$, where $z_i \in \mathbb{R}^d$ constitutes the i -th row of Z and let $w \in \mathbb{R}_{>0}^n$ be a vector of positive weights. Let $I_\beta^+ = \{i \in [n] : w_i z_i^T \beta > 0\}$ and let $I_\beta^- = \{i \in [n] : w_i z_i^T \beta < 0\}$. Let

$$\mu_w(\mathcal{D}) = \sup_{\beta \in \mathbb{R}^d \setminus \{0\}} \frac{\sum_{i \in I_\beta^+} w_i (z_i^T \beta)^2}{\sum_{i \in I_\beta^-} w_i (z_i^T \beta)^2}.$$

We call the dataset \mathcal{D} with weight vector w μ -complex, if there exists a $\mu \in \mathbb{R}$, such that $\mu_w(\mathcal{D}) \leq \mu$.

In order to understand, what exactly this measure does, we have to remember that the parameter vector β that is estimated in the probit model is the orthogonal vector of a hyperplane that separates the space of data points into two partitions. When using the probit model for binary classification, the decision to classify a point as positive or negative is often based on which side of the hyperplane the point is located. If it is possible to perfectly classify each point such that no errors are made, the dataset is called linearly separable. If it isn't linearly separable, there are always two subsets of datapoints: Those points that are correctly classified and those that are not. It turns out, that these two subsets are exactly represented by the sets I_β^+ and I_β^- in the definition of μ . If we assume without losing generality, that for a given $\beta \in \mathbb{R}^d$, we classify a point as positive if $x_i^T \beta < 0$, i.e. it is on the opposite side of the hyperplane where the normal vector β points to, then $z_i^T \beta = -(2y_i - 1)x_i^T \beta > 0$ and I_β^+ contains exactly all the indices of correctly classified points. Conversely, I_β^- then contains all the indices of incorrectly classified points. It is thus easy to see, that if every point is correctly classified, i.e. the dataset is linearly separable, then $I_\beta^- = \emptyset$ and $\mu = \infty$, so the dataset is not μ -complex.

The relationship between linear separability and μ -complexity is even stronger though. In the next theorem, we will show that a finite μ , i.e. the μ -complexity of a dataset, is exactly equivalent to linear separability.

Theorem 3. *Let \mathcal{D} be a d -dimensional dataset of size $|\mathcal{D}| = n$ like in definition 5 (μ -complexity) and let $w \in \mathbb{R}_{>0}^n$ be a vector of positive weights. Then, the dataset \mathcal{D} with weight vector w is μ -complex if and only if \mathcal{D} is not linearly separable.*

Proof. We first prove the " \Rightarrow " direction, i.e. we show that if \mathcal{D} is μ -complex, then it is not linearly separable. We do this by proving the equivalent contraposition that if \mathcal{D} is linearly separable, then it is not μ -complex.

Let $S_0 = \{i \in [n] : y_i = 0\}$ and $S_1 = \{i \in [n] : y_i = 1\}$ like in definition 3 (linear separability). If \mathcal{D} is linearly separable, then there exists a $\beta \in \mathbb{R}^d \setminus \{0\}$, such that

$$\begin{aligned}
& \forall i \in S_0 : x_i^T \beta \geq 0 \quad \text{and} \quad \forall i \in S_1 : x_i^T \beta \leq 0 \\
& \iff \\
& \forall i \in S_0 : (-1)x_i^T \beta \leq 0 \quad \text{and} \quad \forall i \in S_1 : x_i^T \beta \leq 0 \\
& \iff \\
& \forall i \in S_0 : (2y_i - 1)x_i^T \beta \leq 0 \quad \text{and} \quad \forall i \in S_1 : (2y_i - 1)x_i^T \beta \leq 0 \\
& \iff \\
& \forall i \in S_0 : -(2y_i - 1)x_i^T \beta \geq 0 \quad \text{and} \quad \forall i \in S_1 : -(2y_i - 1)x_i^T \beta \geq 0 \\
& \iff \\
& \forall i \in S_0 : z_i^T \beta \geq 0 \quad \text{and} \quad \forall i \in S_1 : z_i^T \beta \geq 0 \\
& \iff \\
& \forall i \in [n] : z_i^T \beta \geq 0 \\
& \iff \\
& I_\beta^- = \{i \in [n] : w_i z_i^T \beta < 0\} = \emptyset \\
& \iff \\
& \sum_{i \in I_\beta^-} w_i (z_i^T \beta)^2 = 0 \\
& \Rightarrow \\
& \mu_w(\mathcal{D}) \geq \frac{\sum_{i \in I_\beta^+} w_i (z_i^T \beta)^2}{\sum_{i \in I_\beta^-} w_i (z_i^T \beta)^2} = \infty,
\end{aligned}$$

which means that \mathcal{D} is not μ -complex.

It now remains to prove the " \Leftarrow " direction, i.e. to show that if \mathcal{D} is not linearly separable, then it is μ -complex. Again, we do this by proving the equivalent contraposition that if \mathcal{D} is not μ -complex, then it is linearly separable.

The first step in order to do so is to show that we can restrict the supremum in $\mu_w(\mathcal{D})$

to finite β with $\|\beta\| = 1$:

$$\begin{aligned}
\mu_w(\mathcal{D}) &= \sup_{\beta \in \mathbb{R}^d \setminus \{0\}} \frac{\sum_{i \in I_\beta^+} w_i (z_i^T \beta)^2}{\sum_{i \in I_\beta^-} w_i (z_i^T \beta)^2} \\
&= \sup_{\beta \in \mathbb{R}^d \setminus \{0\}} \frac{\sum_{i \in I_\beta^+} \frac{1}{\|\beta\|^2} w_i (z_i^T \beta)^2}{\sum_{i \in I_\beta^-} \frac{1}{\|\beta\|^2} w_i (z_i^T \beta)^2} \\
&= \sup_{\beta \in \mathbb{R}^d \setminus \{0\}} \frac{\sum_{i \in I_\beta^+} w_i \left(z_i^T \frac{\beta}{\|\beta\|} \right)^2}{\sum_{i \in I_\beta^-} w_i \left(z_i^T \frac{\beta}{\|\beta\|} \right)^2} \\
&= \sup_{\tilde{\beta} \in \mathbb{R}^d, \|\tilde{\beta}\|=1} \frac{\sum_{i \in I_\beta^+} w_i \left(z_i^T \tilde{\beta} \right)^2}{\sum_{i \in I_\beta^-} w_i \left(z_i^T \tilde{\beta} \right)^2},
\end{aligned}$$

which lets us conclude that even in the supremum, both expressions $\sum_{i \in I_\beta^+} w_i (z_i^T \beta)^2$ and $\sum_{i \in I_\beta^-} w_i (z_i^T \beta)^2$ are finite. This means that if \mathcal{D} is not μ -complex, then the denominator must be zero, i.e. it must hold that there exists a $\beta \in \mathbb{R}^d \setminus \{0\}$ such that

$$\sum_{i \in I_\beta^-} w_i (z_i^T \beta)^2 = 0.$$

From here, we can follow the same chain of equivalences that we showed when proving the " \Rightarrow "-direction of the theorem, which leads us directly to the fact, that \mathcal{D} in this case must be linearly separable, which concludes the proof. \square

Having established the relationship between μ -complexity and linear separability, it directly follows that μ -complexity is also equivalent to the existence and uniqueness of the maximum likelihood estimate of the probit model, that we discussed in section 2.3.

From now on, we will subject our studies of coresets only to those datasets, that are μ -complex, i.e. not linearly separable and with existing and unique maximum likelihood estimate for the probit model.

We conclude this section by proving some simple inequalities regarding μ , which will later be helpful when constructing the coresets.

Lemma 1. *Let \mathcal{D} be a d -dimensional and μ -complex dataset of size $|\mathcal{D}| = n$ with scaled model matrix $Z \in \mathbb{R}^{n \times d}$ and weight vector $w \in \mathbb{R}_{>0}^n$ like in definition 5. The following relationship holds for all $\beta \in \mathbb{R}^d$:*

$$\mu^{-1} \sum_{i \in I_\beta^-} w_i (z_i^T \beta)^2 \leq \sum_{i \in I_\beta^+} w_i (z_i^T \beta)^2 \leq \mu \sum_{i \in I_\beta^-} w_i (z_i^T \beta)^2.$$

Proof. If \mathcal{D} with weights w is μ -complex, then

$$\begin{aligned} & \frac{\sum_{i \in I_{\beta}^+} w_i (z_i^T \beta)^2}{\sum_{i \in I_{\beta}^-} w_i (z_i^T \beta)^2} \leq \mu_w(\mathcal{D}) \leq \mu \\ \iff & \sum_{i \in I_{\beta}^+} w_i (z_i^T \beta)^2 \leq \mu \sum_{i \in I_{\beta}^-} w_i (z_i^T \beta)^2, \end{aligned}$$

which proves the second inequality.

Considering that the labeling of a dataset is arbitrary, i.e. we could always switch the 1 labels for the 0 labels and vice versa (if we flip the sign of β accordingly), the following relationship is true as well:

$$\begin{aligned} & \frac{\sum_{i \in I_{\beta}^-} w_i (z_i^T \beta)^2}{\sum_{i \in I_{\beta}^+} w_i (z_i^T \beta)^2} \leq \mu_w(\mathcal{D}) \leq \mu \\ \iff & \sum_{i \in I_{\beta}^-} w_i (z_i^T \beta)^2 \leq \mu \sum_{i \in I_{\beta}^+} w_i (z_i^T \beta)^2 \\ \iff & \mu^{-1} \sum_{i \in I_{\beta}^-} w_i (z_i^T \beta)^2 \leq \sum_{i \in I_{\beta}^+} w_i (z_i^T \beta)^2, \end{aligned}$$

which proves the first inequality. □

3.2 The Sensitivity Framework

After having imposed some reasonable restrictions on the datasets under study, it is now time to think about how an algorithm that selects a coreset $\mathcal{C} \subseteq \mathcal{D}$ could be constructed. One of the first ideas that come to mind to solve such a problem is the process of random sampling. After all, why don't we just randomly select a subset of points from \mathcal{D} of the desired size? Wouldn't that already solve our problem?

The issue with this approach is that it cannot be guaranteed that such a uniform random sample will yield a coreset, i.e. a subset of \mathcal{D} that helps us to obtain a $(1 \pm \epsilon)$ -approximation of the original loss function. As we will later see in the experiments section, uniform sampling works reasonably well when the data is well behaved, but fails terribly when there are a few very important datapoints that it tends to miss. Intuitively speaking, if there are lots of points in a dataset that don't influence the loss function much, but only a few points that have a big impact on the loss function, uniform sampling fails because it tends to miss the few very important points.

It turns out, that one way to remedy the downsides of uniform sampling is to include a measure of importance in the sampling distribution. Instead of sampling each datapoint with equal probability, why don't we construct our sampling distribution in such a way, that the impact of each point on the loss function is taken into account? This way, more important points would be assigned a higher probability to be sampled and less important points would conversely be assigned a lower sampling probability. This

idea forms the basis of the so called *sensitivity framework*, an algorithmic framework introduced in [Feldman and Langberg, 2011] (see also [Feldman et al., 2020]), that aims to find coresets by randomly sampling points proportional to their importance for the loss function.

In the sensitivity framework, the importance of a point in a dataset can be thought of the maximum proportion of the loss function that it can take up in the worst case. To formalize this intuition, the sensitivity framework shifts the representation of a dataset as a collection of points towards the representation as a collection of *functions*, where each function represents the loss of a point. To explain what that means, consider the dataset $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$ with scaled model matrix Z , i.e. the rows of Z are given by vectors $z_i = -(2y_i - 1)x_i$ and the loss function $f(\beta) = \sum_{i=1}^n w_i g(z_i^T \beta)$, where w_1, \dots, w_n are positive weights. From now on, we will assign to each point in the dataset the function $g_i(\beta) = g(z_i^T \beta)$, that represents its individual contribution to the overall loss function. This way, we can equivalently represent the dataset \mathcal{D} as the set of functions $F = \{g_1, \dots, g_n\}$.

Having made this conceptual change of representing a dataset as a set of functions, we can now use this new representation to formalize the concept of importance for each point, according to which we later want to sample. As already hinted at, the importance of a point will be the maximum share of the loss function, that the loss of the specific point will take up in the worst case. This worst case importance is also called the *sensitivity* of a point, and it was first introduced in [Langberg and Schulman, 2010]. A formal definition of this concept, which forms the basis of the sensitivity framework, is given below.

Definition 6 ([Langberg and Schulman, 2010]). *Let $F = \{g_1, \dots, g_n\}$ be a set of functions, $g_i : \mathbb{R}^d \rightarrow \mathbb{R}_{\geq 0}$, $i \in [n]$ and let $w \in \mathbb{R}_{>0}^n$ be a vector of positive weights. The sensitivity of g_i for $f_w(\beta) = \sum_{i=1}^n w_i g_i(\beta)$ is defined as*

$$\varsigma_i = \sup_{\beta \in \mathbb{R}^d, f_w(\beta) > 0} \frac{w_i g_i(\beta)}{f_w(\beta)}.$$

The total sensitivity, i.e. the sum of the sensitivities is $\mathfrak{S} = \sum_{i=1}^n \varsigma_i$.

The idea behind the sensitivity framework is to sample from a distribution where the sampling probabilities are determined by the sensitivities, i.e. the worst case importance of each point in the dataset. But there is one issue remaining: The true sensitivities $\varsigma_1, \dots, \varsigma_n$ are unknown, and as pointed out in [Braverman et al., 2020], their computation requires solving the original optimization problem, which we wanted to avoid in the first place. Luckily, there exists a simple workaround: Instead of sampling proportionally to the true sensitivities, we can instead also sample proportionally to upper bounds s_i , where $s_i \geq \varsigma_i$, that are potentially easier to compute.

There is one caveat though, that we have to take into account: We will later see, that in the sensitivity framework, the size of the coreset that we obtain through sampling proportionally to upper bounds of the sensitivities, is directly influenced by the sum of

said bounds: $S = \sum_{i=1}^n s_i$. It follows, that we have to find bounds that are as tight as possible, so that our coreset won't be unnecessarily large.

The way in which the authors of [Feldman and Langberg, 2011] were able to show that sampling proportionally to upper bounds of the sensitivities can lead to provably small coresets, was to relate the concept of sensitivities to the theory of so-called *range spaces* and the *VC-dimension* (see for example [Kearns and Vazirani, 1994] for an introduction of the VC-dimension). We introduce these concepts in the following definitions, because they will also turn out to be of crucial importance to the size of our coresets.

Definition 7 ([Feldman and Langberg, 2011]). *A range space is a pair $\mathfrak{R} = (F, \text{ranges})$, where F is a set and ranges is a family (set) of subsets of F .*

Definition 8 ([Feldman and Langberg, 2011]). *The VC-dimension $\Delta(\mathfrak{R})$ of a range space $\mathfrak{R} = (F, \text{ranges})$ is the size $|G|$ of the largest subset $G \subseteq F$ such that*

$$|\{G \cap R \mid R \in \text{ranges}\}| = 2^{|G|},$$

i.e. G is shattered by ranges .

The most important part to understand about the two preceding definitions is the concept of *shattering*. Given a set F , a set of subsets of F called *ranges* and a subset $G \subseteq F$, what does it mean if G is shattered by *ranges*? Definition 8 says, that the set of intersections between G and *ranges* must be equal to $2^{|G|}$, which is exactly the size of the set of all subsets of G . It follows, that for G to be shattered by *ranges*, every subset of G must appear in (be a subset of) at least one of the sets in *ranges*. From this it also immediately follows, that if G is shattered by *ranges*, then every subset of G is also shattered by *ranges*. Thus, the VC-dimension of the range space that is given by F and *ranges* is simply the size of the largest subset $G \subseteq F$, such that every subset of G appears in at least one element of *ranges*.

Instead of dealing with arbitrary sets like in definition 7 and definition 8, we are specifically dealing with a set of functions $F = \{g_1, \dots, g_n\}$, that represents our dataset. In the next definition, we will see, how such a set of functions can be used to *induce* a range space, which will be of crucial importance when limiting the size of our coresets later on.

Definition 9 ([Feldman and Langberg, 2011]). *Let F be a finite set of functions mapping from \mathbb{R}^d to $\mathbb{R}_{\geq 0}$. For every $\beta \in \mathbb{R}^d$ and $r \geq 0$, let*

$$\text{range}(F, \beta, r) = \{f \in F \mid f(\beta) \geq r\}$$

and let

$$\text{ranges}(F) = \{\text{range}(F, \beta, r) \mid \beta \in \mathbb{R}^d, r \geq 0\}.$$

Then we call $\mathfrak{R}_F := (F, \text{ranges}(F))$ the range space induced by F .

To understand what it means for a set of functions F to induce a range space, consider that we can always partition the set of functions into two disjoint subsets by choosing

a specific β and applying a threshold r : Every element in F with a value of $f(\beta) \geq r$ goes to one subset, the remainder goes to the other subset. Considering that our functions actually represent datapoints, each value of β and r give one way to partition the datapoints into two disjoint subsets. Thus, the set $\text{ranges}(F)$ in definition 9 can be thought of the set of all possible partitions of functions (datapoints), that can be obtained for all possible $\beta \in \mathbb{R}^d$ and $r \geq 0$ and the VC-dimension of the induced range space is the size of the largest set of datapoints, that can be arbitrarily partitioned by different values of β and r .

We are now ready for the full theorem that relates sensitivity sampling to the theory of range spaces and the VC-dimension, which forms the core of the sensitivity framework. Its original version goes back to [Feldman and Langberg, 2011] and it was further improved in [Braverman et al., 2020]. The version we are presenting here is a slightly adapted variant introduced in [Feldman et al., 2020]:

Theorem 4 ([Feldman et al., 2020]). *Let $F = \{g_1, \dots, g_n\}$ be a finite set of functions mapping from \mathbb{R}^d to $\mathbb{R}_{\geq 0}$. Let $w \in \mathbb{R}_{>0}^n$ be a vector of positive weights. Let $\epsilon, \delta \in (0, \frac{1}{2})$. Let $s_i \geq \varsigma_i$ be upper bounds of the sensitivities and let $S = \sum_{i=1}^n s_i$. Given s_i , one can compute in time $O(|F|)$ a set $R \subseteq F$ of*

$$|R| \in O\left(\frac{S}{\epsilon^2} \left(\Delta \log S + \log\left(\frac{1}{\delta}\right)\right)\right)$$

weighted functions, such that with probability $1 - \delta$ we have for all $\beta \in \mathbb{R}^d$ simultaneously

$$(1 - \epsilon) \sum_{g_i \in F} w_i g_i(\beta) \leq \sum_{g_i \in R} u_i g_i(\beta) \leq (1 + \epsilon) \sum_{g_i \in F} w_i g_i(\beta).$$

Each element of R is sampled independently with probability $p_j = \frac{s_j}{S}$ from F , $u_i = \frac{Sw_j}{s_j|R|}$ denotes the weight of a function $g_i \in R$ that corresponds to $g_j \in F$ and Δ is an upper bound on the VC-dimension of the range space \mathfrak{R}_{F^} induced by F^* , where F^* is the set of functions $g_i \in F$ scaled by $\frac{Sw_i}{s_i|R|}$, i.e. $F^* = \left\{ \frac{Sw_i}{s_i|R|} g_i(\beta) \mid i \in [n] \right\}$.*

The set of functions $R \subseteq F$ with weights u is of course a representation of the coreset that we are interested in. As theorem 4 tells us, the size of this coreset depends on both, the sum of the sensitivity bounds as well as the VC-dimension of the range space induced by F^* , a reweighted version of F .

Another thing to note is that theorem 4 introduces a failure probability δ , which also influences the size of the coreset. The reason why we need this new parameter becomes clear when considering, that the coreset is selected by random sampling, i.e. the coreset-construction process is probabilistic and probabilistic processes can fail. This failure probability is reflected by the parameter δ .

Equipped with theorem 4, we now have a clear roadmap to follow in search of our goal of developing the coreset construction algorithm. We know, that we have to find small upper bounds on the sensitivities of our function set $F = \{g_1, \dots, g_n\}$ for the purpose of random sampling and at the same time, we also know that we have to control the VC-dimension of the range space of F^* . Thus, bounding the sensitivities as well as bounding the VC-dimension are the main challenges of the next section.

3.3 Constructing the Coreset

Without any point of reference, the task of finding tight and efficiently computable upper bounds on the sensitivities seems rather challenging. How are we supposed to find those ominous bounds and on top of that, how can we make sure that their sum will be small? It helps to remind ourselves of the original problem that the sensitivities were supposed to solve. Instead of sampling every point with equal probability, the sensitivities are used to include the importance of each point into the sampling process. But what if similar importance distributions already existed? Could it be possible to choose an existing importance distribution and relate it to the concept of sensitivities in order to obtain upper bounds?

It turns out, that there is one importance sampling distribution that is particularly helpful in the context of our problem: The so-called statistical leverage scores (see for example [Drineas et al., 2012]). For a dataset $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$ with model matrix X , the statistical leverage score of the i -th observation is given by $\ell_i = x_i^T (X^T X)^{-1} x_i$. Intuitively speaking, ℓ_i is a measure of the uniqueness of a given observation x_i . If there aren't many data points close to x_i , i.e. we can consider x_i to be unique, then the leverage score of x_i is high. On the other hand, if there are a lot of other data points close to x_i , i.e. x_i is not unique, then the leverage score is low. When using leverage scores as a sampling distribution, we give a higher weight to points that are unique and different from the other points, compared to points that are surrounded by a lot of similar other points.

An alternative way of defining the leverage scores, which will be particularly important for us, is to specify them as the squared row norms of an orthonormal basis of the model matrix $X \in \mathbb{R}^{n \times d}$. Such an orthonormal basis can for example be obtained by conducting a so-called QR -decomposition of X (see for example [Golub and van Loan, 2013]), where $X = QR$ is factorized into an orthonormal matrix $Q \in \mathbb{R}^{n \times d}$ and an upper triangular matrix $R \in \mathbb{R}^{d \times d}$.

Our goal in this section is to adapt the leverage scores in such a way, that we can obtain tight upper bounds on the sensitivities. But we are not quite ready for that yet. Before we can get there, we first have to focus our attention back to the probit loss function $g(x)$, because it turns out that in order to bound the sensitivities by using the leverage scores, we first have to cover some important properties of $g(x)$.

3.3.1 A closer examination of the probit loss

In order to find bounds on the sensitivities, we also need bounds on the probit loss $g(x) = \ln\left(\frac{1}{1-\Phi(x)}\right)$. The first bound that we will derive holds for all $x \geq 0$ and shows that $g(x)$ grows at least like a quadratic function. Next, we will show that for all $x \geq 2$, $g(x)$ is upper bounded by a quadratic function, and thus $g(x)$ asymptotically grows like a quadratic function. Both of these bounds will turn out to be helpful later on. They are proven in lemma 2 and lemma 3.

Lemma 2. Let $g(x) = \ln \left(\frac{1}{1-\Phi(x)} \right)$. Then, for all $x \geq 0$, it holds that:

$$\frac{1}{2}x^2 \leq g(x).$$

Proof. We first show the claim for all $x \geq 1$, by using the following inequality:

$$\begin{aligned} \Phi(-x) &= \frac{1}{\sqrt{2\pi}} \int_{-\infty}^{-x} \exp\left(-\frac{1}{2}t^2\right) dt \\ &\leq \frac{1}{\sqrt{2\pi}} \int_{-\infty}^{-x} -t \exp\left(-\frac{1}{2}t^2\right) dt \\ &= \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{1}{2}x^2\right) \\ &\leq \exp\left(-\frac{1}{2}x^2\right). \end{aligned}$$

In the next step, we use this inequality to show that for $x \geq 1$:

$$\begin{aligned} e^{g(x)} &= e^{\ln\left(\frac{1}{1-\Phi(x)}\right)} = \frac{1}{\Phi(-x)} \geq e^{\frac{1}{2}x^2} \\ &\iff \\ g(x) &\geq \frac{1}{2}x^2, \end{aligned}$$

which proves the bound for $x \geq 1$.

Let us now turn to the case when $0 \leq x \leq 1$. Both $g(x)$ and $\frac{1}{2}x^2$ are monotonically increasing and continuous functions for $0 \leq x \leq 1$. Making use of the fact that $g(0) > \frac{1}{2}$, it follows for all $0 \leq x \leq 1$, that

$$g(x) \geq g(0) > \frac{1}{2} = \max_{0 \leq x \leq 1} \frac{1}{2}x^2 \geq \frac{1}{2}x^2,$$

which concludes the proof. □

Lemma 3. Let $g(x) = \ln \left(\frac{1}{1-\Phi(x)} \right)$. Then, for all $x \geq 2$, it holds that:

$$g(x) \leq x^2.$$

Proof. In [Gordon, 1941], it was shown that the following inequality holds for all $x \geq 0$:

$$\Phi(-x) \geq \frac{1}{\sqrt{2\pi}} \frac{x}{x^2 + 1} e^{-\frac{1}{2}x^2}.$$

We can use this inequality to establish that for all $x \geq 2$ it holds that:

$$\begin{aligned}
e^{x^2} \cdot \Phi(-x) &\geq e^{x^2} \frac{1}{\sqrt{2\pi}} \frac{x}{x^2 + 1} e^{-\frac{1}{2}x^2} \\
&= e^{\frac{1}{2}x^2} \frac{1}{\sqrt{2\pi}} \frac{x}{x^2 + 1} \\
&= e^{\frac{1}{2}x^2} \frac{1}{\frac{4}{3}(x^2 + 1)} \frac{\frac{4}{3}x}{\sqrt{2\pi}} \\
&\geq \frac{e^{\frac{1}{2}x^2}}{\frac{4}{3}(x^2 + 1)} \\
&\geq \frac{e^{\frac{1}{2}x^2}}{e^{\frac{1}{2}x^2}} \\
&= 1 \\
&\iff \\
e^{x^2} &\geq \frac{1}{1 - \Phi(x)} \\
&\iff \\
x^2 &\geq \ln \left(\frac{1}{1 - \Phi(x)} \right) = g(x),
\end{aligned}$$

which completes the proof. \square

3.3.2 Finding the sensitivity bounds

Having successfully established some quadratic bounds on the probit loss $g(x)$, we can now turn our attention back to the task of finding upper bounds on the sensitivities. As already mentioned, we will be using the statistical leverage scores in order to do so.

To this end, we will show that the function set $F = \{g_1, \dots, g_n\}$, which represents our dataset, can be partitioned into two classes of functions for which we can find upper bounds on the sensitivities. The first class will contain all points, where the loss function surpasses a specific threshold. It turns out, that for a given β , the threshold $z_i^T \beta \geq 2$ is a suitable candidate. In the next lemma, we show how we can relate the loss function of all the points in this class to the statistical leverage scores, while also incorporating μ into the upper bound. To do this, we use the notion of leverage scores as the squared row norms of an orthogonal basis, that can be obtained for example by conducting a QR -factorization.

Lemma 4. *Let \mathcal{D} be a d -dimensional and μ -complex dataset of size $|\mathcal{D}| = n$ with scaled model matrix $Z \in \mathbb{R}^{n \times d}$ and let $w \in \mathbb{R}_{>0}^n$ be a vector of positive weights. Let $F = \{g_1, \dots, g_n\}$ be a set of functions with $g_i(\beta) = g(z_i^T \beta)$ and let $f_w(\beta) = \sum_{i=1}^n w_i g_i(\beta)$. Then, it holds that*

$$w_j g_j(\beta) \leq 2 \|U_j\|_2^2 (1 + \mu) f_w(\beta) \quad \forall j \in \{i \in [n] : z_i^T \beta \geq 2\},$$

where $U \in \mathbb{R}^{n \times d}$ is an orthonormal basis for the columnspace of $\sqrt{D_w}Z$ and $\sqrt{D_w} \in \mathbb{R}^{n \times n}$ is a diagonal matrix, where the i -th diagonal element is equal to $\sqrt{w_i}$ and $U_j \in \mathbb{R}^d$ is the j -th row of U .

Proof. Let $\sqrt{D_w}Z = UR$, where U is an orthonormal basis for the columnspace of $\sqrt{D_w}Z$. Then, for all $j \in \{i \in [n] : z_i^T \beta \geq 2\}$:

$$w_j g_j(\beta) = w_j g(z_j^T \beta) = w_j g\left(\frac{\sqrt{w_j} z_j^T \beta}{\sqrt{w_j}}\right) = w_j g\left(\frac{U_j^T R \beta}{\sqrt{w_j}}\right) \leq w_j g\left(\frac{\|U_j\|_2 \|R \beta\|_2}{\sqrt{w_j}}\right),$$

where $U_j \in \mathbb{R}^d$ is the vector that constitutes the j 'th row of U and the inequality is true due to the *Cauchy-Schwarz inequality*. We continue the proof as follows:

$$\begin{aligned} w_j g\left(\frac{\|U_j\|_2 \|R \beta\|_2}{\sqrt{w_j}}\right) &= w_j g\left(\frac{\|U_j\|_2 \|UR \beta\|_2}{\sqrt{w_j}}\right) \\ &= w_j g\left(\frac{\|U_j\|_2 \|\sqrt{D_w} Z \beta\|_2}{\sqrt{w_j}}\right) \\ &\leq \|U_j\|_2^2 \|\sqrt{D_w} Z \beta\|_2^2 \\ &= \|U_j\|_2^2 \sum_{i=1}^n w_i (z_i^T \beta)^2. \end{aligned}$$

Here, the first equality follows from the fact that U is orthonormal, i.e. multiplying by U doesn't change the norm of a vector. The inequality follows from the bound $g(x) \leq x^2$ that holds for all $x \geq 2$, which was shown in lemma 3.

Now, let $I_\beta^+ = \{i \in [n] : w_i z_i^T \beta > 0\}$ and let $I_\beta^- = \{i \in [n] : w_i z_i^T \beta < 0\}$ like in definition 5, the definition of μ -complexity. We continue the proof by making use of the relationship that was shown in lemma 1:

$$\begin{aligned} \|U_j\|_2^2 \sum_{i=1}^n w_i (z_i^T \beta)^2 &= \|U_j\|_2^2 \left(\sum_{i \in I_\beta^+} w_i (z_i^T \beta)^2 + \sum_{i \in I_\beta^-} w_i (z_i^T \beta)^2 \right) \\ &\leq \|U_j\|_2^2 \left(\sum_{i \in I_\beta^+} w_i (z_i^T \beta)^2 + \mu \sum_{i \in I_\beta^+} w_i (z_i^T \beta)^2 \right) \\ &= \|U_j\|_2^2 (1 + \mu) \sum_{i \in I_\beta^+} w_i (z_i^T \beta)^2 \\ &\leq 2 \|U_j\|_2^2 (1 + \mu) \sum_{i \in I_\beta^+} w_i g(z_i^T \beta), \end{aligned}$$

where the last inequality follows from the bound $g(x) \geq \frac{1}{2}x^2$, that holds for all $x \geq 0$, which we proved in lemma 2.

From here, we can use the fact that g is a strictly positive function to complete the proof:

$$2\|U_j\|_2^2(1+\mu) \sum_{i \in I_\beta^+} w_i g(z_i^T \beta) \leq 2\|U_j\|_2^2(1+\mu) \sum_{i=1}^n w_i g(z_i^T \beta) = 2\|U_j\|_2^2(1+\mu) f_w(\beta)$$

□

In the next lemma, we turn to the remaining class of points where $z_i^T \beta \leq 2$, and show how the sensitivity of these points can be bounded by a constant value that only depends on μ and the weight of the given point.

Lemma 5. *Let \mathcal{D} be a d -dimensional and μ -complex dataset of size $|\mathcal{D}| = n$ with scaled model matrix $Z \in \mathbb{R}^{n \times d}$ and let $w \in \mathbb{R}_{>0}^n$ be a vector of positive weights. Let $F = \{g_1, \dots, g_n\}$ be a set of functions with $g_i(\beta) = g(z_i^T \beta)$ and let $f_w(\beta) = \sum_{i=1}^n w_i g_i(\beta)$. Then, it holds that*

$$w_j g_j(\beta) \leq \frac{w_j}{\mathcal{W}} (80 + 16\mu) f_w(\beta) \quad \forall j \in \{i \in [n] : z_i^T \beta \leq 2\},$$

where $\mathcal{W} = \sum_{i=1}^n w_i$ is the sum of all weights.

Proof. We first start by noting that $g(-1) > \frac{1}{10}$ and that $g(2) < 4$. Now, we partition the indices into two sets as follows:

$$\begin{aligned} K_\beta^- &= \{i \in [n] \mid z_i^T \beta \leq -1\} \\ K_\beta^+ &= \{i \in [n] \mid z_i^T \beta > -1\}. \end{aligned}$$

In the case that $\sum_{j \in K_\beta^+} w_j \geq \frac{1}{2} \mathcal{W}$, the following relationship holds:

$$f_w(\beta) = \sum_{i=1}^n w_i g(z_i^T \beta) \geq \sum_{i \in K_\beta^+} w_i g(z_i^T \beta) \geq \frac{\sum_{i \in K_\beta^+} w_i}{10} \geq \frac{\mathcal{W}}{20} = \frac{\mathcal{W}}{20w_j} w_j \geq \frac{\mathcal{W}}{80w_j} w_j g(z_j^T \beta),$$

where $j \in \{i \in [n] : z_i^T \beta \leq 2\}$. Thus, we have in this case:

$$w_j g(z_j^T \beta) \leq \frac{80w_j}{\mathcal{W}} f_w(\beta).$$

If on the other hand $\sum_{j \in K_\beta^-} w_j \geq \frac{1}{2} \mathcal{W}$, we have that

$$f_w(\beta) = \sum_{i=1}^n w_i g(z_i^T \beta) \geq \sum_{i \in I_\beta^+} w_i g(z_i^T \beta) \geq \frac{1}{2} \sum_{i \in I_\beta^+} w_i (z_i^T \beta)^2 \geq \frac{1}{2\mu} \sum_{i \in I_\beta^-} w_i (z_i^T \beta)^2,$$

where $I_\beta^+ = \{i \in [n] : w_i z_i^T \beta > 0\}$ and $I_\beta^- = \{i \in [n] : w_i z_i^T \beta < 0\}$ like in definition 5 (μ -complexity). The second inequality is true due to the lower bound $g(x) \geq \frac{1}{2}x^2$ that

holds for all $x \geq 0$ (see lemma 2) and the third inequality is true due to a property of μ that was proved in lemma 1.

We continue the proof as follows:

$$\frac{1}{2\mu} \sum_{i \in I_\beta^-} w_i (z_i^T \beta)^2 \geq \frac{1}{2\mu} \sum_{i \in K_\beta^-} w_i (z_i^T \beta)^2 \geq \frac{1}{2\mu} \sum_{i \in K_\beta^-} w_i \geq \frac{\mathcal{W}}{4\mu} \geq \frac{\mathcal{W}}{16\mu w_j} w_j g(z_j^T \beta),$$

which leads us to the upper bound for the second case:

$$w_j g(z_j^T \beta) \leq \frac{16\mu w_j}{\mathcal{W}} f_w(\beta).$$

We can conclude the proof by adding both upper bounds:

$$w_j g_j(\beta) = w_j g(z_j^T \beta) \leq \frac{80w_j}{\mathcal{W}} f_w(\beta) + \frac{16\mu w_j}{\mathcal{W}} f_w(\beta) = \frac{w_j}{\mathcal{W}} (80 + 16\mu) f_w(\beta).$$

□

It is now time to use the results from lemma 4 and lemma 5 to derive an upper bound on the sensitivities. Because we showed how the dataset can be partitioned in "high loss points" and "low loss points" for every given β and how these two classes of points can both be upper bounded, it simply suffices to add those bounds together to bound the sensitivities for any given β . As a final step, we only have to show that the total sum of the sensitivities remains small. We do both in the following lemma.

Lemma 6. *Let \mathcal{D} be a d -dimensional and μ -complex dataset of size $|\mathcal{D}| = n$ with scaled model matrix $Z \in \mathbb{R}^{n \times d}$, let $w \in \mathbb{R}_{>0}^n$ be a vector of positive weights and let $U \in \mathbb{R}^{n \times d}$ be an orthonormal basis for the columnspace of $\sqrt{D_w}Z$. Let $F = \{g_1, \dots, g_n\}$ be a set of functions with $g_i(\beta) = g(z_i^T \beta)$ and let $f_w(\beta) = \sum_{i=1}^n w_i g_i(\beta)$. Then, the sensitivity ς_i of g_i (see definition 6) is upper bounded by*

$$\varsigma_i \leq s_i = (80 + 16\mu) (\|U_i\|_2^2 + \frac{w_i}{\mathcal{W}}),$$

and the total sensitivity is bounded by

$$\mathfrak{S} = \sum_{i=1}^n \varsigma_i \leq 192\mu d.$$

Proof. We can use the bounds that we derived in lemma 4 and lemma 5 to bound the sensitivities:

$$\begin{aligned} \varsigma_i &= \sup_{\beta \in \mathbb{R}^d, f_w(\beta) > 0} \frac{w_i g(z_i^T \beta)}{f_w(\beta)} \\ &\leq \sup_{\beta \in \mathbb{R}^d, f_w(\beta) > 0} \frac{2\|U_i\|_2^2(1 + \mu)f_w(\beta) + \frac{w_i}{\mathcal{W}}(80 + 16\mu)f_w(\beta)}{f_w(\beta)} \\ &= 2\|U_i\|_2^2(1 + \mu) + \frac{w_i}{\mathcal{W}}(80 + 16\mu) \\ &\leq \|U_i\|_2^2(80 + 16\mu) + \frac{w_i}{\mathcal{W}}(80 + 16\mu) \\ &= (80 + 16\mu) (\|U_i\|_2^2 + \frac{w_i}{\mathcal{W}}), \end{aligned}$$

which completes the first part of the proof. For the next part, we use that U is an orthonormal matrix. The Frobenius norm $\|U\|_F$ (see for example [Golub and van Loan, 2013]) of an orthonormal matrix is equal to \sqrt{d} , as can easily be verified:

$$\|U\|_F = \sqrt{\sum_{k=1}^d \sum_{l=1}^n |u_{lk}|^2} = \sqrt{\sum_{k=1}^d 1} = \sqrt{d},$$

where the second equality follows from the fact that the columns of U have unit norm due to its orthonormality. We can now conclude the proof as follows:

$$\begin{aligned} \mathfrak{S} &= \sum_{i=1}^n \varsigma_i \leq (80 + 16\mu) \sum_{i=1}^n \|U_i\|_2^2 + \frac{w_i}{\mathcal{W}} \\ &= (80 + 16\mu)(\|U\|_F^2 + 1) \\ &= (80 + 16\mu)(d + 1) \\ &\leq 96\mu(d + 1) \\ &\leq 192\mu d. \end{aligned}$$

□

We now have successfully completed the first task on our list of developing a coresets construction algorithm by using the sensitivity framework. Not only did we manage to derive upper bounds on the sensitivities of the function set $F = \{g_1, \dots, g_n\}$ that represents our dataset by using the statistical leverage scores, but we also showed that the sum of those bounds is in $O(\mu d)$, even independent of the total number of datapoints n ! The final step, before putting everything together, is now to find an upper bound of the VC-dimension of the range-space induced by F^* .

3.3.3 Bounding the VC-dimension

Remember, that in the core theorem of the sensitivity framework (see theorem 4), the function-set that induces our range space of interest is defined as

$$F^* = \left\{ \frac{S w_i}{s_i |R|} g_i(\beta) \mid i \in [n] \right\},$$

where s_i are upper bounds on the sensitivities, S is the sum of these upper bounds, $|R|$ is the size of the sample and w_i are the initial weights. We are thus dealing with a set of weighted probit loss functions.

In order to approach the complex problem of bounding the VC-dimension of a set of arbitrarily weighted probit loss functions, we first deal with the slightly simpler problem that arises when the weights only consist of a single positive constant, i.e. we are looking at the set $\mathcal{F}^c = \{c g_i(\beta) \mid i \in [n]\}$. The authors of [Huggins et al., 2016] showed, that in the case of the logistic loss, it is possible to relate the VC-dimension of the range space induced by \mathcal{F}^c to the VC-dimension of the affine hyperplane classifier, which the

authors of [Kearns and Vazirani, 1994] showed to be bounded by $d+1$. It turns out, that a similar case can be made for the probit loss, which we demonstrate in the following lemma.

Lemma 7 (cf. [Huggins et al., 2016]). *Let $Z \in \mathbb{R}^{n \times d}$, let $z_i \in \mathbb{R}^d$ be the i -th row of Z and let $c \in \mathbb{R}_{>0}$. Let $F = \{g_1, \dots, g_n\}$ be a set of functions with $g_i(\beta) = g(z_i^T \beta)$, where $g(x) = \ln\left(\frac{1}{1-\Phi(x)}\right)$ is the probit loss. The VC-dimension of the range space induced by*

$$\mathcal{F}^c = \{cg_i(\beta) \mid i \in [n]\}$$

is bounded by $\Delta(\mathfrak{R}_{\mathcal{F}^c}) \leq d+1$.

Proof. We start by noting that for all $G \subseteq \mathcal{F}^c$ we have

$$|\{G \cap R \mid R \in \text{ranges}(\mathcal{F}^c)\}| = |\{\text{range}(G, \beta, r) \mid \beta \in \mathbb{R}^d, r \geq 0\}|.$$

Since g is invertible and monotone, we have for all $\beta \in \mathbb{R}^d$ and $r \geq 0$ that

$$\begin{aligned} \text{range}(G, \beta, r) &= \{g_i \in G \mid g_i(\beta) \geq r\} \\ &= \{g_i \in G \mid cg_i(\beta) \geq r\} \\ &= \left\{g_i \in G \mid z_i^T \beta \geq g^{-1}\left(\frac{r}{c}\right)\right\}. \end{aligned}$$

Note, that $\{g_i \in G \mid z_i^T \beta \geq g^{-1}\left(\frac{r}{c}\right)\}$ corresponds to the positively classified points of the affine hyperplane classifier $x \mapsto \text{sign}(x^T \beta - g^{-1}\left(\frac{r}{c}\right))$. We thus have for all $G \subseteq \mathcal{F}^c$, that

$$|\{G \cap R \mid R \in \text{ranges}(\mathcal{F}^c)\}| = |\{\{g_i \in G \mid z_i^T \beta - s \geq 0\} \mid \beta \in \mathbb{R}^d, s \in \mathbb{R}\}|.$$

As shown in [Kearns and Vazirani, 1994], the VC-dimension of the set of affine hyperplane classifiers is $d+1$, so it follows that $\Delta(\mathfrak{R}_{\mathcal{F}^c}) \leq d+1$, which concludes the proof. \square

In the next step, we will generalize the class \mathcal{F}^c of constantly weighted probit loss functions to the class $\mathcal{F}^w = \{w_i g_i(\beta) \mid i \in [n]\}$ of arbitrarily weighted probit loss functions for a weight vector $w \in \mathbb{R}_{>0}^n$, which also includes F^* , our class of interest. The authors of [Munteanu et al., 2018] presented a proof that shows how the VC-dimension of the range space induced by such a set can be bounded by $t \cdot (d+1)$ in the case of logistic regression, where $t \in \mathbb{N}$ is the number of distinct weights in the vector w . We follow a similar path and adapt their argument to the context of probit regression in the following lemma.

Lemma 8 (cf. [Munteanu et al., 2018]). *Let $Z \in \mathbb{R}^{n \times d}$, let $z_i \in \mathbb{R}^d$ be the i -th row of Z and let $w \in \mathbb{R}_{>0}^n$ be a vector of positive weights, where $w_i \in \{v_1, \dots, v_t\}$ for all $i \in [n]$. Let $F = \{g_1, \dots, g_n\}$ be a set of functions with $g_i(\beta) = g(z_i^T \beta)$. The VC-dimension of the range space induced by*

$$\mathcal{F}^w = \{w_i g_i(\beta) \mid i \in [n]\}$$

is bounded by $\Delta(\mathfrak{R}_{\mathcal{F}^w}) \leq t \cdot (d+1)$.

Proof. We start by partitioning the functions in \mathcal{F}^w into t disjoint classes

$$F_j = \{w_i g(z_i \beta) \in \mathcal{F}^w \mid w_i = v_j\}, \quad j \in [t].$$

The functions in each of these classes have an equal weight, which means that by lemma 7, each of their induced range spaces has a VC-dimension of at most $d + 1$.

For the sake of contradiction, assume that $\Delta(\mathfrak{R}_{\mathcal{F}^w}) > t \cdot (d + 1)$ and let G be the corresponding set of size $|G| > t \cdot (d + 1)$ that is shattered by $\text{ranges}(\mathcal{F}^w)$. Since the sets F_j are disjoint, each intersection $F_j \cap G$ must be shattered by $\text{ranges}(F_j)$ as well. Further, at least one of the intersections must have at minimum $\frac{|G|}{t}$ elements, which means that for at least one $j \in [t]$ it holds that $|F_j \cap G| \geq \frac{|G|}{t} > \frac{t \cdot (d+1)}{t} = d + 1$. This is a contradiction to lemma 7, which concludes the proof. \square

We have now found a way to bound the VC-dimension of the range space induced by F^* in the number of distinct weights, i.e. the number of distinct values of $\frac{S w_i}{s_i |R|}$. But there is one important issue that remains to be dealt with: In the general case, we don't know the number of distinct values of $\frac{S w_i}{s_i |R|}$, and it is even reasonable to assume, that this value can be equal to the total number of datapoints, n . This would be a problem for us though, because the core theorem of the sensitivity framework (theorem 4) tells us, that the size of our coresset will depend linearly on the VC-dimension of the range space induced by F^* . If this VC-dimension is in $O(n)$, our coresets won't be small. It follows, that we have to find a way to work around that problem in order to obtain small coresets.

3.3.4 A simple two-pass algorithm for small coresets

It now remains to solve the final challenge: How can we limit the number of distinct weights in F^* in order to limit the VC-dimension of the range space induced by F^* and obtain small coresets? The authors of [Munteanu et al., 2018] were facing a similar situation, and it turns out that they managed to come up with a very clever idea: Their approach to the problem is to slightly increase the upper bounds on the sensitivities s_i to a new value $s'_i \geq s_i$, such that the fraction $\frac{s'_i}{w_i}$ is exactly a power of two. This way, the total sum of the sensitivities $S' = \sum_{i=1}^n s'_i$ is still under control, because as a worst case bound we have that $S' \leq 2S$, which doesn't influence the big-o notation and we still have $S' \in O(\mu d)$. The big advantage of this approach is, that we now have a way to bound the number of distinct values of $\frac{S' w_i}{s'_i |R|}$ in a term that is logarithmic in n . To see why this is the case, we first derive two simple inequalities. The first one goes like this and holds for all $i \in [n]$:

$$\frac{s'_i}{w_i} \leq \frac{2s_i}{w_i} = \frac{2(80 + 16\mu)(\|U_i\|_2^2 + \frac{w_i}{\mathcal{W}})}{w_i} \leq \frac{192\mu(\|U_i\|_2^2 + \frac{w_i}{\mathcal{W}})}{w_i} \leq \frac{384\mu}{w_i} \leq \frac{384\mu}{w_{\min}},$$

where w_{\min} is the minimum weight and \mathcal{W} is the sum of all weights. The third inequality holds, because it is always true that $\mu \geq 1$ and the fourth inequality is true because one property of the statistical leverage scores is that $\|U_i\|_2^2 \leq 1$.

Next, we show how $\frac{s'_i}{w_i}$ can be lower-bounded for all $i \in [n]$:

$$\frac{s'_i}{w_i} \geq \frac{s_i}{w_i} \geq \sup_{\beta \in \mathbb{R}^d} \frac{g_i(\beta)}{\sum_{i=1}^n w_i g_i(\beta)} \stackrel{\beta=0}{\geq} \frac{1}{\sum_{i=1}^n w_i} \geq \frac{1}{nw_{\max}},$$

where w_{\max} is the maximum weight. Putting both of the inequalities together, we thus have that

$$\frac{1}{nw_{\max}} \leq \frac{s'_i}{w_i} \leq \frac{384\mu}{w_{\min}}.$$

We know, that the values of $\frac{s'_i}{w_i}$ are exactly powers of two, so we can compute the amount of possible distinct values of $\frac{s'_i}{w_i}$, which we call t , like this:

$$t \leq \log_2 \left(\frac{384\mu}{w_{\min}} \right) - \log_2 \left(\frac{1}{nw_{\max}} \right) = \log_2 \left(384\mu n \frac{w_{\max}}{w_{\min}} \right) \in O(\log_2(\mu n \omega)),$$

where $\omega = \frac{w_{\max}}{w_{\min}}$. Thus, when sampling according to s'_i , our function class of interest becomes

$$F^{*'} = \left\{ \frac{S'w_i}{s'_i|R|} g_i(\beta) \mid i \in [n] \right\},$$

and the weights $\frac{S'w_i}{s'_i|R|}$ can assume only $O(\log_2(\mu n \omega))$ distinct values. Plugging this into lemma 8, we get that the VC-dimension of the range space induced by $F^{*'}$ is upper bounded by a term in $O(d \log_2(\mu n \omega))$.

Theorem 5. *Let \mathcal{D} be a d -dimensional and μ -complex dataset of size $|\mathcal{D}| = n$ with scaled model matrix $Z \in \mathbb{R}^{n \times d}$, let $w \in \mathbb{R}_{>0}^n$ be a vector of positive weights, with $\omega = \frac{w_{\max}}{w_{\min}}$ being the ratio of the largest and smallest weight, $\mathcal{W} = \sum_{i=1}^n w_i$ being the sum of all weights, and let $U \in \mathbb{R}^{n \times d}$ be an orthonormal basis for the columnspace of $\sqrt{D_w}Z$, where $U_i \in \mathbb{R}^d$ is the vector that constitutes the i -th row of U . Let $\epsilon \in (0, \frac{1}{2})$.*

If $\mathcal{C} \subseteq \mathcal{D}$ is a subset of \mathcal{D} of size $|\mathcal{C}| = k$, that was obtained by independently sampling

$$k \in O \left(\frac{\mu d^2}{\epsilon^2} \log(\omega n) \log(\mu d) \right)$$

elements from \mathcal{D} proportional to

$$q_i = \min \left\{ 2^l \mid l \in \mathbb{Z}, 2^l \geq \|U_i\|_2^2 + \frac{w_i}{\mathcal{W}} \right\},$$

i.e. with sampling probability $p_i = \frac{q_i}{\sum_{i=1}^n q_i}$ for all $i \in [n]$ and $u \in \mathbb{R}_{>0}^k$ is a new weight vector, where $u_j = \frac{w_i \sum_{l=1}^n p_l}{k p_i}$ is the new weight for an element in \mathcal{C} that corresponds to the i -th element of \mathcal{D} , then with probability $1 - n^{-c}$, \mathcal{C} with weights u is a $(1 \pm \epsilon)$ -coreset of \mathcal{D} for probit regression for any absolute constant $c > 1$.

Proof. $S \leq 2 \cdot 192\mu d$, $\Delta = d \log(\omega n)$, $\delta = \log^{-c}(n)$.

$$\begin{aligned} k &\in O\left(\frac{S}{\epsilon^2} \left(\Delta \log S + \log\left(\frac{1}{\delta}\right)\right)\right) \\ &\subseteq O\left(\frac{\mu d}{\epsilon^2} (d \log(\omega n) \log(\mu d) + \log(n^c))\right) \\ &\subseteq O\left(\frac{\mu d^2}{\epsilon^2} \log(\omega n) \log(\mu d)\right) \end{aligned}$$

□

4 Data Streams

Content.

5 Experiments

Content.

6 Concluding Remarks

Content.

7 Notes

7.1 VC Dimension

An alternative approach is to write down the VC dimension by using an instance space and a concept class as given in [Kearns and Vazirani, 1994].

Lemma 9. *Let $X = \{x_1, \dots, x_n\} \subset \mathbb{R}^d \times \mathbb{R}_{>0}$ be the instance space consisting of n points with their last coordinate being positive. The concept class of interest, \mathcal{C} over X , is given as follows:*

$$\mathcal{C} = \{ \{x \in X : f_{\beta,r}(x) \geq 0\} \mid \beta \in \mathbb{R}^d, r \geq 0 \},$$

with

$$f_{\beta,r}(x) = x_{d+1} \cdot g\left(\sum_{i=1}^d x_i \beta_i\right) - r$$

and

$$g(x) = -\log \Phi(-x).$$

The VC dimension of \mathcal{C} is equal to the VC dimension of the range space induced by $\mathcal{F}_{probit}^w = \{w_i g(z_i \beta) \mid i \in [n]\}$, $Z \in \mathbb{R}^{n \times d}$, $w \in \mathbb{R}_{>0}^n$.

There are a few different strategies that can be used to find an upper bound on the VC dimension of \mathcal{C} , as shown by the following lemmas. The first one is a simple upper bound for finite concept classes:

Lemma 10. *Let X be an instance space and \mathcal{C} be a concept class over X . If the cardinality of \mathcal{C} can be bounded by m , i.e. $|\mathcal{C}| \leq m$, then $VCdim(\mathcal{C}) \leq \log(m)$.*

The next lemma partitions the concept class into smaller classes, for each of which the VC dimension can be bounded:

Lemma 11. *Let X be an instance space and \mathcal{C} be a concept class over X . Let $\mathcal{C}_1, \dots, \mathcal{C}_k$ be a partition of \mathcal{C} into k disjoint subsets, i.e. $\mathcal{C} = \bigcup_{i=1}^k \mathcal{C}_i$ and $\mathcal{C}_i \cap \mathcal{C}_j = \emptyset \forall i \neq j$. Then, $VCdim(\mathcal{C}) \leq \sum_{i=1}^k VCdim(\mathcal{C}_i)$.*

Proof. For the sake of contradiction, assume there was a set $S \subseteq X$ of size $|S| > \sum_{i=1}^k VCdim(\mathcal{C}_i)$ that is shattered by \mathcal{C} . If S is shattered by \mathcal{C} , every subset of S must also be shattered by \mathcal{C} . Consider the intersections $T_i = \bigcup_{c \in \mathcal{C}_i} S \cap c$. Every T_i is a subset of S and $S = \bigcup_{i=1}^k T_i$. Since S is shattered by \mathcal{C} , every T_i must be shattered by \mathcal{C}_i . We assumed that $|S| > \sum_{i=1}^k VCdim(\mathcal{C}_i)$. It follows that there exists a T_j with $|T_j| > VCdim(\mathcal{C}_j)$. Since T_j is also shattered by \mathcal{C}_j , this is a contradiction, which concludes the proof. \square

A result in [Linial et al., 1991] suggests an even smaller upper bound:

Lemma 12 ([Linial et al., 1991]). *Let X be an instance space and \mathcal{C} be a concept class over X . Let $\mathcal{C} = \bigcup_{i=1}^k \mathcal{C}_i$ and $VCdim(\mathcal{C}_i) \leq m$. If k is bounded by a polynomial function of m , then $VCdim(\mathcal{C}) \leq 3m$.*

Instead of partitioning the concept class, we could also partition the instance space and obtain a similar bound:

Lemma 13. *Let X be an instance space and \mathcal{C} be a concept class over X . Let X_1, \dots, X_k be a partition of X into k disjoint subsets, i.e. $X = \bigcup_{i=1}^k X_i$ and $X_i \cap X_j = \emptyset \ \forall i \neq j$. Let $\mathcal{C}_i = \{X_i \cap c \mid c \in \mathcal{C}\}$ be a concept class over X_i for all $i \in [k]$. Then, $VCdim(\mathcal{C}) \leq \sum_{i=1}^k VCdim(\mathcal{C}_i)$.*

Proof. Again, assume there existed a set $S \subseteq X$ of size $|S| > \sum_{i=1}^k VCdim(\mathcal{C}_i)$ that is shattered by \mathcal{C} . S can be partitioned into disjoint subsets $T_i = S \cap X_i$, with $\bigcup_{i=1}^k T_i = S$. Every T_i must be shattered by \mathcal{C}_i . Since we assumed that $|S| > \sum_{i=1}^k VCdim(\mathcal{C}_i)$, there exists a T_j with $|T_j| > VCdim(\mathcal{C}_j)$ which is also shattered by \mathcal{C}_j . This contradiction concludes the proof. \square

7.2 New idea for VC dimension proof

Lemma 14. *Let*

$$h_{\beta,r}(x) = \begin{cases} 1 & \text{if } x_{d+1} \cdot g\left(\sum_{i=1}^d x_i \beta_i\right) - r \geq 0 \\ 0 & \text{else} \end{cases}$$

Be a function from \mathbb{R}^{d+1} to $\{0, 1\}$ with parameters $\beta \in \mathbb{R}^d$ and $r \in \mathbb{R}_{\geq 0}$ with

$$g(x) = \log\left(\frac{1}{1 - \Phi(x)}\right),$$

$$\Phi(x) = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^x e^{-\frac{1}{2}z^2} dz.$$

Let

$$H = \{x \mapsto h_{\beta,r}(x) \mid \beta \in \mathbb{R}^d, r \in \mathbb{R}_{\geq 0}\}$$

be the hypothesis class determined by h . Then, the VC dimension of H is ...

Proof. Let $S = \sum_{i=1}^d x_i \beta_i$. We show that h can be computed in t steps as follows:

$$\begin{aligned} & x_{d+1} \cdot g(S) - r \geq 0 \\ \iff & \log\left(\frac{1}{1 - \Phi(S)}\right) \geq \frac{r}{x_{d+1}} \\ \iff & \frac{1}{1 - \Phi(S)} \geq \exp\left(\frac{r}{x_{d+1}}\right) \\ \iff & 1 - \Phi(S) \leq \exp\left(-\frac{r}{x_{d+1}}\right) \\ \iff & \Phi(S) \geq 1 - \exp\left(-\frac{r}{x_{d+1}}\right) \\ \iff & \frac{1}{\sqrt{2\pi}} \int_{-\infty}^S e^{-\frac{1}{2}z^2} dz \geq 1 - \exp\left(-\frac{r}{x_{d+1}}\right) \\ \iff & \int_{-\infty}^S e^{-\frac{1}{2}z^2} dz \geq \sqrt{2\pi} \left(1 - \exp\left(-\frac{r}{x_{d+1}}\right)\right) \end{aligned}$$

□

7.3 Online Leverage Scores

The leverage scores of a matrix $A \in \mathbb{R}^{n \times d}$ are given by $l_i = a_i^T (A^T A)^{-1} a_i$ [Cohen et al., 2020]. According to [Cohen et al., 2020], we can obtain overestimates of these scores by using only a subset of the rows in A to compute them.

Let A_j be a matrix that contains only the first j rows of A . It follows that the estimated leverage score $\tilde{l}_j = a_j^T (A_j^T A_j)^{-1} a_j$ is an overestimate of l_j . In a recent paper by [Chhaya et al., 2020], it was shown that the sum of these overestimates can be bounded regardless of how the rows in A are ordered:

Lemma 15 ([Chhaya et al., 2020]).

$$\sum_{i=1}^n \tilde{l}_i \in O(d + d \log \|A\| - \min_{i \in [n]} \|a_i\|)$$

Next, we show how a simple algorithm that computes \tilde{l}_j in an online manner (passing row by row over the data stream) can be constructed requiring only $\mathcal{O}(d^2)$ of working memory. The idea is to only keep the matrix $A_j^T A_j \in \mathbb{R}^{d \times d}$ in memory and update it for every new row a_{j+1} using a rank one update $A_{j+1}^T A_{j+1} = A_j^T A_j + a_{j+1} \cdot a_{j+1}^T$. See [Golub and van Loan, 2013] for more on matrix multiplication using outer products. The algorithm is given in algorithm 2.

Algorithm 2: Online Leverage Scores

Input: Matrix $A \in \mathbb{R}^{n \times d}$

Output: Online leverage scores \tilde{l}_i for all $i \in [n]$

- 1 Initialize $M_0 = 0^{d \times d}$
 - 2 **foreach** $a_i := i$ 'th row vector of A , $a_i \in \mathbb{R}^d$ **do**
 - 3 $M_i = M_{i-1} + a_i \cdot a_i^T$
 - 4 $\tilde{l}_i = a_i^T M_i^\dagger a_i$
 - 5 **return** $\tilde{l}_i, i \in [n]$
-

References

- [Agresti, 2015] Agresti, A. (2015). *Foundations of Linear and Generalized Linear Models*. Wiley.
- [Albert and Chib, 1993] Albert, J. H. and Chib, S. (1993). Bayesian analysis of binary and polychotomous response data. *Journal of the American Statistical Association*, 88(422):669–679.
- [Braverman et al., 2020] Braverman, V., Feldman, D., Lang, H., Statman, A., and Zhou, S. (2020). New frameworks for offline and streaming coresets constructions.
- [Chhaya et al., 2020] Chhaya, R., Choudhari, J., Dasgupta, A., and Shit, S. (2020). Streaming coresets for symmetric tensor factorization. *CoRR*, abs/2006.01225.
- [Cohen et al., 2020] Cohen, M. B., Musco, C., and Pachocki, J. (2020). Online row sampling. *Theory of Computing*, 16(15):1–25.
- [Demidenko, 2001] Demidenko, E. (2001). Computational aspects of probit model. *Mathematical Communications*, 6:233–247.
- [Drineas et al., 2012] Drineas, P., Magdon-Ismail, M., Mahoney, M. W., and Woodruff, D. P. (2012). Fast approximation of matrix coherence and statistical leverage. *Journal of Machine Learning Research*, 13(111):3475–3506.
- [Fahrmeir et al., 2013] Fahrmeir, L., Kneib, T., Lang, S., and Marx, B. D. (2013). *Regression*. Springer-Verlag Berlin Heidelberg.
- [Feldman and Langberg, 2011] Feldman, D. and Langberg, M. (2011). A unified framework for approximating and clustering data. In *Proceedings of the Forty-Third Annual ACM Symposium on Theory of Computing*, STOC ’11, page 569–578, New York, NY, USA. Association for Computing Machinery.
- [Feldman et al., 2020] Feldman, D., Schmidt, M., and Sohler, C. (2020). Turning big data into tiny data: Constant-size coresets for k -means, pca, and projective clustering. *SIAM Journal on Computing*, 49(3):601–657.
- [Gelfand and Smith, 1990] Gelfand, A. E. and Smith, A. F. M. (1990). Sampling-based approaches to calculating marginal densities. *Journal of the American Statistical Association*, 85(410):398–409.
- [Gelman et al., 2013] Gelman, A., Carlin, J., Stern, H., Dunson, D., Vehtari, A., and Rubin, D. (2013). *Bayesian Data Analysis*. Chapman and Hall/CRC, 3 edition.
- [Golub and van Loan, 2013] Golub, G. H. and van Loan, C. F. (2013). *Matrix Computations*. JHU Press, fourth edition.

- [Gordon, 1941] Gordon, R. D. (1941). Values of mills’ ratio of area to bounding ordinate and of the normal probability integral for large values of the argument. *The Annals of Mathematical Statistics*, 12(3):364–366.
- [Haberman, 1974] Haberman, S. J. (1974). *The Analysis of Frequency Data*. The University of Chicago Press.
- [Huggins et al., 2016] Huggins, J. H., Campbell, T., and Broderick, T. (2016). Coresets for scalable bayesian logistic regression. In *Proceedings of the 30th International Conference on Neural Information Processing Systems*, NIPS’16, page 4087–4095, Red Hook, NY, USA. Curran Associates Inc.
- [Kearns and Vazirani, 1994] Kearns, M. J. and Vazirani, U. V. (1994). *An Introduction to Computational Learning Theory*. MIT Press.
- [Kremer et al., 1999] Kremer, I., Nisan, N., and Ron, D. (1999). On randomized one-round communication complexity. *Computational Complexity*, 8(1):21–49.
- [Langberg and Schulman, 2010] Langberg, M. and Schulman, L. J. (2010). Universal ϵ -approximators for integrals. In *Proceedings of the Twenty-First Annual ACM-SIAM Symposium on Discrete Algorithms*, SODA ’10, pages 598–607.
- [Lesaffre and Kaufmann, 1992] Lesaffre, E. and Kaufmann, H. (1992). Existence and uniqueness of the maximum likelihood estimator for a multivariate probit model. *Journal of the American Statistical Association*, 87:805–811.
- [Linial et al., 1991] Linial, N., Mansour, Y., and Rivest, R. L. (1991). Results on learnability and the vapnik-chervonenkis dimension. *Information and Computation*, 90(1):33–49.
- [McCullagh and Nelder, 1989] McCullagh, P. and Nelder, J. A. (1989). *Generalized Linear Models*. Chapman and Hall/CRC.
- [Munteanu and Schwiegelshohn, 2018] Munteanu, A. and Schwiegelshohn, C. (2018). Coresets-methods and history: A theoreticians design pattern for approximation and streaming algorithms. *KI - Künstliche Intelligenz*, 32(1):37–53.
- [Munteanu et al., 2018] Munteanu, A., Schwiegelshohn, C., Sohler, C., and Woodruff, D. P. (2018). On coresets for logistic regression. In *Advances in Neural Information Processing Systems 31, (NeurIPS)*, pages 6562–6571.
- [Nocedal and Wright, 2006] Nocedal, J. and Wright, S. (2006). *Numerical Optimization*. Springer-Verlag New York, 2 edition.
- [Tanner and Wong, 1987] Tanner, M. A. and Wong, W. H. (1987). The calculation of posterior distributions by data augmentation. *Journal of the American Statistical Association*, 82(398):528–540.

- [Vaart, 1998] Vaart, A. W. v. d. (1998). *Asymptotic Statistics*. Cambridge Series in Statistical and Probabilistic Mathematics. Cambridge University Press.
- [Wedderburn, 1976] Wedderburn, R. W. M. (1976). On the existence and uniqueness of the maximum likelihood estimates for certain generalized linear models. *Biometrika*, 63(1):27–32.